

# Visa\_Approval\_Prediction

May 30, 2024

- 1 Objective : Visa Approval Classification
- 2 Exploratory Data Analysis (EDA) - Python
- 3 Insights - Patterns
- 4 Classification (Using the ML)



## 5 1. Load Python Modules

```
[1]: # Use Python's import statement to load modules
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns

import warnings
warnings.filterwarnings('ignore', category=FutureWarning)
```

```

from tabulate import tabulate

from sklearn.base import BaseEstimator, TransformerMixin
from sklearn.pipeline import Pipeline
from sklearn.preprocessing import LabelEncoder

from sklearn.preprocessing import StandardScaler
from sklearn.preprocessing import MinMaxScaler
from sklearn.model_selection import train_test_split

from sklearn.linear_model import LogisticRegression

from sklearn.naive_bayes import CategoricalNB
from sklearn.naive_bayes import GaussianNB

from sklearn.svm import SVC

from sklearn.tree import DecisionTreeClassifier
from sklearn.ensemble import RandomForestClassifier

from sklearn.ensemble import AdaBoostClassifier
from sklearn.tree import DecisionTreeClassifier

from sklearn.ensemble import GradientBoostingClassifier
from xgboost import XGBClassifier

from sklearn.metrics import accuracy_score
from sklearn.metrics import classification_report
from sklearn.metrics import confusion_matrix
from imblearn.over_sampling import SMOTE

```

## 6 2. Read the Dataset from CSV file - Using Pandas

```

[2]: file_path=r"Visa_Predaction_Dataset.csv"
visa_df=pd.read_csv(file_path)
visa_df

```

```

[2]:
   case_id continent education_of_employee has_job_experience \
0    EZYV01      Asia      High School                N
1    EZYV02      Asia      Master's                Y
2    EZYV03      Asia      Bachelor's                N
3    EZYV04      Asia      Bachelor's                N
4    EZYV05  Africa      Master's                Y
...      ...      ...      ...      ...

```

25475	EZYV25476	Asia	Bachelor's	Y
25476	EZYV25477	Asia	High School	Y
25477	EZYV25478	Asia	Master's	Y
25478	EZYV25479	Asia	Master's	Y
25479	EZYV25480	Asia	Bachelor's	Y

	requires_job_training	no_of_employees	yr_of_estab	\
0	N	14513	2007	
1	N	2412	2002	
2	Y	44444	2008	
3	N	98	1897	
4	N	1082	2005	
...	...	...	...	
25475	Y	2601	2008	
25476	N	3274	2006	
25477	N	1121	1910	
25478	Y	1918	1887	
25479	N	3195	1960	

	region_of_employment	prevailing_wage	unit_of_wage	full_time_position	\
0	West	592.2029	Hour	Y	
1	Northeast	83425.6500	Year	Y	
2	West	122996.8600	Year	Y	
3	West	83434.0300	Year	Y	
4	South	149907.3900	Year	Y	
...	...	...	...	...	
25475	South	77092.5700	Year	Y	
25476	Northeast	279174.7900	Year	Y	
25477	South	146298.8500	Year	N	
25478	West	86154.7700	Year	Y	
25479	Midwest	70876.9100	Year	Y	

	case_status
0	Denied
1	Certified
2	Denied
3	Denied
4	Certified
...	...
25475	Certified
25476	Certified
25477	Certified
25478	Certified
25479	Certified

[25480 rows x 12 columns]

```
[3]: #drop - sensitive - non imp columns for data analysis
print(visa_df["case_id"].nunique())
visa_df.drop("case_id",axis=1,inplace=True)
# print columns names
print(visa_df.columns)
```

25480

```
Index(['continent', 'education_of_employee', 'has_job_experience',
      'requires_job_training', 'no_of_employees', 'yr_of_estab',
      'region_of_employment', 'prevailing_wage', 'unit_of_wage',
      'full_time_position', 'case_status'],
      dtype='object')
```

## 7 3. Basic Inspection on given dataset

```
[4]: def basic_inspection_dataset(table):
      """Generates a basic inspection dataset from the given table."""

      print("top 5 rows - using head")
      print(table.head())
      print()

      print("bottom 5 rows using tail")
      print(table.tail())
      print()

      print("numbers of samples and columns")
      print(table.shape)
      print()

      print("numbers of samples ")
      print(len(table))
      print()

      print("numbers of entries in the data frame")
      print(table.size)
      print()

      print("Columns Names")
      print(table.columns)
      print()

      print("Columns dtypes")
      print(table.dtypes)
      print()

      print("Dataframe info")
```

```

print(table.info())
print()

print()
print("check the missing value in each column")
print(table.isnull().sum())

print()
print("check the missing value in each column")
print(table.isna().sum())

print()
print("table describe")
print(table.describe())

basic_inspection_dataset(visa_df)

```

top 5 rows - using head

	continent	education_of_employee	has_job_experience	requires_job_training	\
0	Asia	High School	N	N	
1	Asia	Master's	Y	N	
2	Asia	Bachelor's	N	Y	
3	Asia	Bachelor's	N	N	
4	Africa	Master's	Y	N	

	no_of_employees	yr_of_estab	region_of_employment	prevailing_wage	\
0	14513	2007	West	592.2029	
1	2412	2002	Northeast	83425.6500	
2	44444	2008	West	122996.8600	
3	98	1897	West	83434.0300	
4	1082	2005	South	149907.3900	

	unit_of_wage	full_time_position	case_status
0	Hour	Y	Denied
1	Year	Y	Certified
2	Year	Y	Denied
3	Year	Y	Denied
4	Year	Y	Certified

bottom 5 rows using tail

	continent	education_of_employee	has_job_experience	\
25475	Asia	Bachelor's	Y	
25476	Asia	High School	Y	
25477	Asia	Master's	Y	
25478	Asia	Master's	Y	
25479	Asia	Bachelor's	Y	

	requires_job_training	no_of_employees	yr_of_estab	\
25475	Y	2601	2008	
25476	N	3274	2006	
25477	N	1121	1910	
25478	Y	1918	1887	
25479	N	3195	1960	

	region_of_employment	prevailing_wage	unit_of_wage	full_time_position	\
25475	South	77092.57	Year		Y
25476	Northeast	279174.79	Year		Y
25477	South	146298.85	Year		N
25478	West	86154.77	Year		Y
25479	Midwest	70876.91	Year		Y

	case_status
25475	Certified
25476	Certified
25477	Certified
25478	Certified
25479	Certified

numbers of samples and columns  
(25480, 11)

numbers of samples  
25480

numbers of entries in the data frame  
280280

Columns Names

```
Index(['continent', 'education_of_employee', 'has_job_experience',
      'requires_job_training', 'no_of_employees', 'yr_of_estab',
      'region_of_employment', 'prevailing_wage', 'unit_of_wage',
      'full_time_position', 'case_status'],
      dtype='object')
```

Columns dtypes

continent	object
education_of_employee	object
has_job_experience	object
requires_job_training	object
no_of_employees	int64
yr_of_estab	int64
region_of_employment	object
prevailing_wage	float64
unit_of_wage	object
full_time_position	object

```
case_status          object
dtype: object
```

Dataframe info

```
<class 'pandas.core.frame.DataFrame'>
```

```
RangeIndex: 25480 entries, 0 to 25479
```

```
Data columns (total 11 columns):
```

#	Column	Non-Null Count	Dtype
0	continent	25480 non-null	object
1	education_of_employee	25480 non-null	object
2	has_job_experience	25480 non-null	object
3	requires_job_training	25480 non-null	object
4	no_of_employees	25480 non-null	int64
5	yr_of_estab	25480 non-null	int64
6	region_of_employment	25480 non-null	object
7	prevailing_wage	25480 non-null	float64
8	unit_of_wage	25480 non-null	object
9	full_time_position	25480 non-null	object
10	case_status	25480 non-null	object

```
dtypes: float64(1), int64(2), object(8)
```

```
memory usage: 2.1+ MB
```

```
None
```

check the missing value in each column

```
continent          0
education_of_employee  0
has_job_experience  0
requires_job_training  0
no_of_employees    0
yr_of_estab        0
region_of_employment  0
prevailing_wage    0
unit_of_wage       0
full_time_position  0
case_status        0
dtype: int64
```

check the missing value in each column

```
continent          0
education_of_employee  0
has_job_experience  0
requires_job_training  0
no_of_employees    0
yr_of_estab        0
region_of_employment  0
prevailing_wage    0
```

```
unit_of_wage          0
full_time_position    0
case_status           0
dtype: int64
```

```
table describe
      no_of_employees  yr_of_estab  prevailing_wage
count      25480.000000  25480.000000    25480.000000
mean        5667.043210   1979.409929    74455.814592
std         22877.928848    42.366929    52815.942327
min          -26.000000   1800.000000      2.136700
25%         1022.000000   1976.000000    34015.480000
50%         2109.000000   1997.000000    70308.210000
75%         3504.000000   2005.000000   107735.512500
max        602069.000000  2016.000000   319210.270000
```

### 7.0.1 Observations - dataset

- Have 25480 Sample with Variables 12
- There is no null values in the dataset

#### Categorical Variables:

- case\_id
- continent
- education\_of\_employee
- has\_job\_experience
- requires\_job\_training
- unit\_of\_wage
- full\_time\_position
- case\_status
- yr\_of\_estab
- region\_of\_employment

#### Numerical Variables:

- prevailing\_wage
- no\_of\_employees

## 8 4. Handling Missing Values - Categorical - Variables

```
[5]: # check for missing values - for confirmation
visa_df.isnull().sum()
```

```
[5]: continent          0
education_of_employee  0
has_job_experience     0
requires_job_training  0
```



```

no_of_employees      0
yr_of_estab          0
region_of_employment 0
prevailing_wage       0
unit_of_wage         0
full_time_position    0
case_status           0
dtype: int64

```

## 9 5. Categorical- UniVariate - Analysis -Using Pipeline

```

[6]: class BarPieChartTransformer(BaseEstimator, TransformerMixin):
    def __init__(self):
        pass

    def fit(self, X, y=None):
        return self

    def transform(self, X):
        df=X.copy()
        # get cat columns
        cat_cols = df.select_dtypes(include='object').columns
        for cat_name in cat_cols:
            value_counts = df[cat_name].value_counts().reset_index()
            # Rename the columns
            value_counts.columns = ['Class', 'Frequency']

            # Print the result as a table
            print(f"{cat_name} frequency table")
            print(tabulate(value_counts, headers='keys', tablefmt='pretty'))

            # Calculate relative frequency
            total_count = value_counts['Frequency'].sum()
            value_counts['Relative Frequency %'] =
↳round((value_counts['Frequency'] / total_count)*100,2)

            # Print the result as a table
            print(f"{cat_name} Relative frequency table")
            print(tabulate(value_counts, headers='keys', tablefmt='pretty'))

            # Extract the values and index from value counts
            value_counts = df[cat_name].value_counts()
            values = value_counts.values
            labels = value_counts.index

        fig, axs = plt.subplots(1, 2, figsize=(12, 6)) # 1 row, 2 columns

```

```

# Create a bar graph
axs[0].bar(labels, values)
axs[0].set_title(f'Frequency of {cat_name}')
axs[0].set_xlabel('Categories') # Set x-label
axs[0].set_ylabel('Count')      # Set y-label

axs[1].pie(value_counts.values, labels=value_counts.index,
↳autopct='%1.1f%%', startangle=140)
axs[1].set_title(f'Relative Frequency of {cat_name}')
plt.tight_layout()
# Show the plot
plt.show()

```

```

[7]: pipeline_cat_var = Pipeline([
      ('cat_univariate_analysis', BarPieChartTransformer())
    ])

# Fit and transform your data using the pipeline
processed_data = pipeline_cat_var.fit_transform(visa_df)


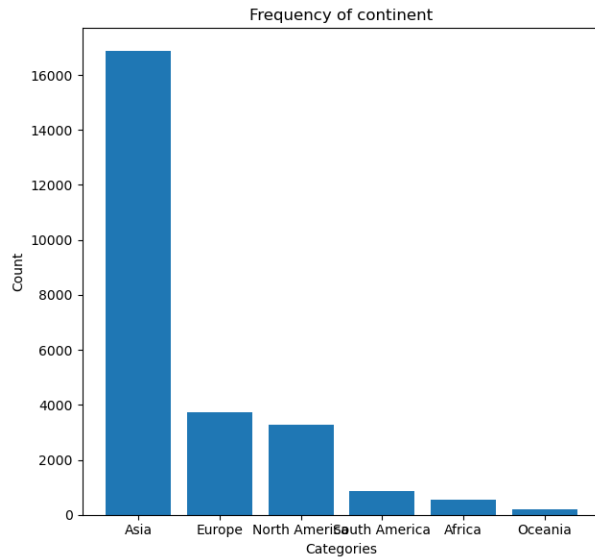
```

continent frequency table

	Class	Frequency
0	Asia	16861
1	Europe	3732
2	North America	3292
3	South America	852
4	Africa	551
5	Oceania	192

continent Relative frequency table

	Class	Frequency	Relative Frequency %
0	Asia	16861	66.17
1	Europe	3732	14.65
2	North America	3292	12.92
3	South America	852	3.34
4	Africa	551	2.16
5	Oceania	192	0.75

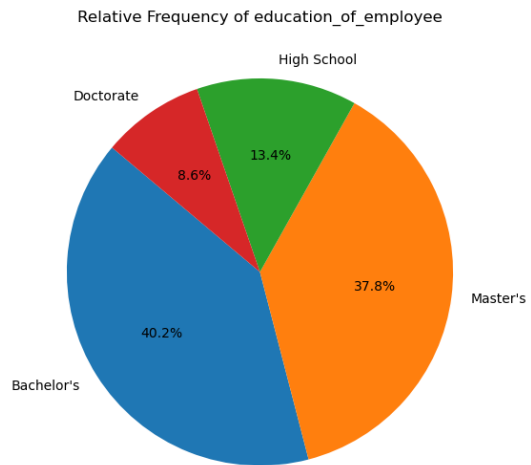
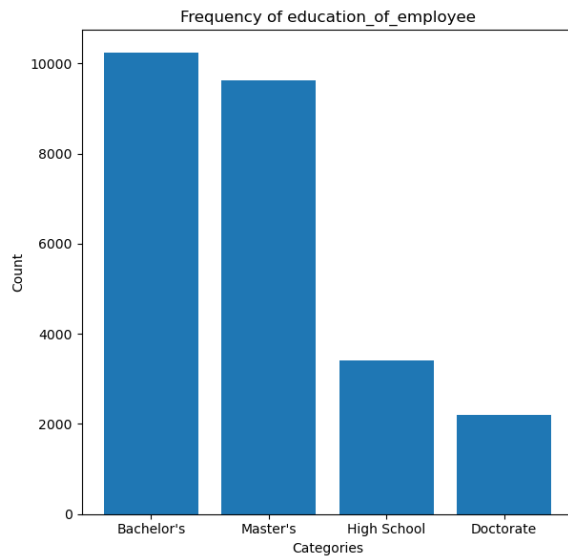


education\_of\_employee frequency table

	Class	Frequency
0	Bachelor's	10234
1	Master's	9634
2	High School	3420
3	Doctorate	2192

education\_of\_employee Relative frequency table

	Class	Frequency	Relative Frequency %
0	Bachelor's	10234	40.16
1	Master's	9634	37.81
2	High School	3420	13.42
3	Doctorate	2192	8.6


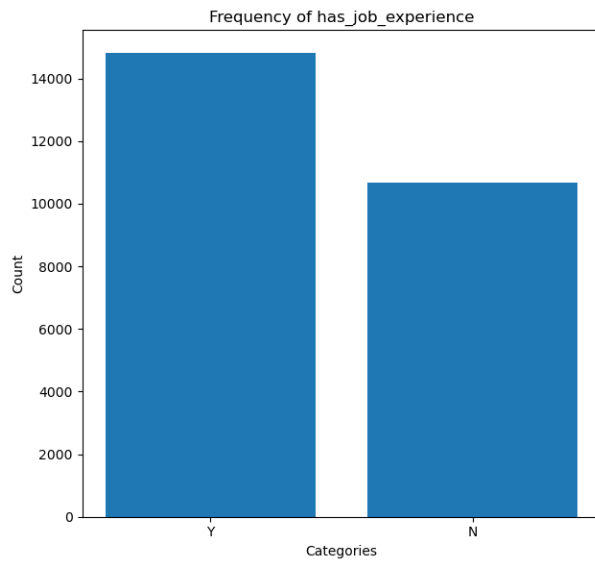


has\_job\_experience frequency table

	Class	Frequency
0	Y	14802
1	N	10678

has\_job\_experience Relative frequency table

	Class	Frequency	Relative Frequency %
0	Y	14802	58.09
1	N	10678	41.91





requires\_job\_training frequency table

	Class	Frequency
0	N	22525
1	Y	2955

requires\_job\_training Relative frequency table

	Class	Frequency	Relative Frequency %
0	N	22525	88.4
1	Y	2955	11.6

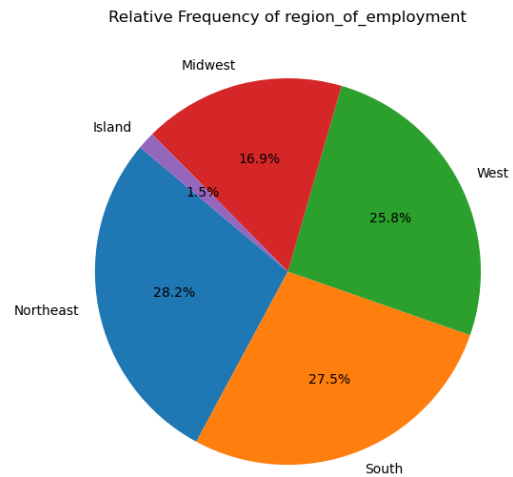
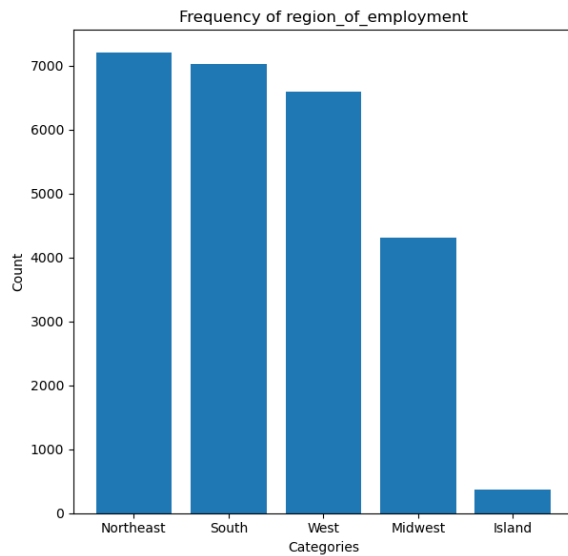


region\_of\_employment frequency table

	Class	Frequency
0	Northeast	7195
1	South	7017
2	West	6586
3	Midwest	4307
4	Island	375

region\_of\_employment Relative frequency table

	Class	Frequency	Relative Frequency %
0	Northeast	7195	28.24
1	South	7017	27.54
2	West	6586	25.85
3	Midwest	4307	16.9
4	Island	375	1.47

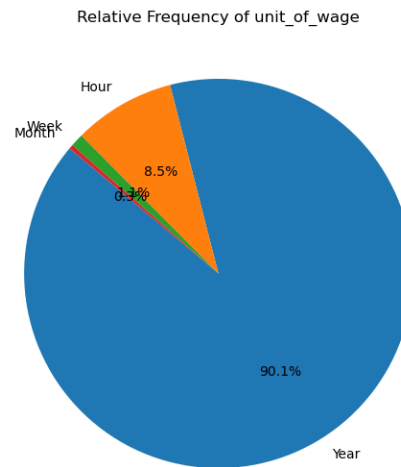
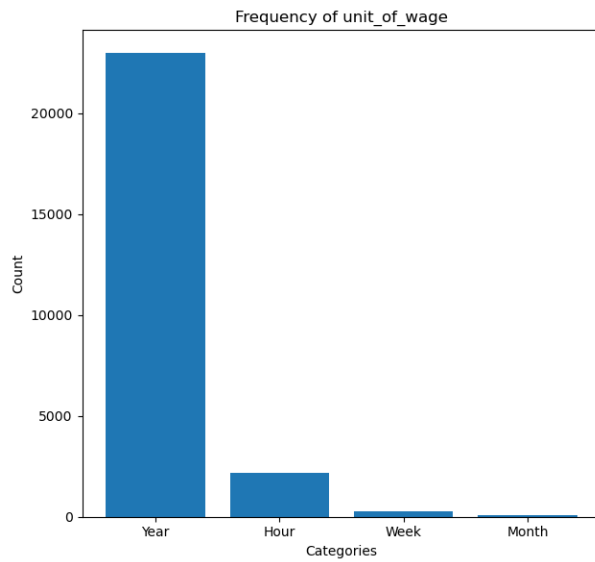


unit\_of\_wage frequency table

	Class	Frequency
0	Year	22962
1	Hour	2157
2	Week	272
3	Month	89

unit\_of\_wage Relative frequency table

	Class	Frequency	Relative Frequency %
0	Year	22962	90.12
1	Hour	2157	8.47
2	Week	272	1.07
3	Month	89	0.35




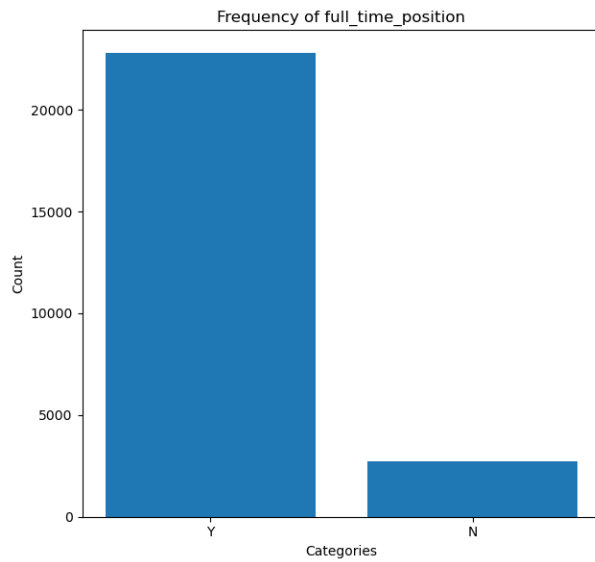
full\_time\_position frequency table

	Class	Frequency
0	Y	22773
1	N	2707

full\_time\_position Relative frequency table

	Class	Frequency	Relative Frequency %
0	Y	22773	89.38
1	N	2707	10.62



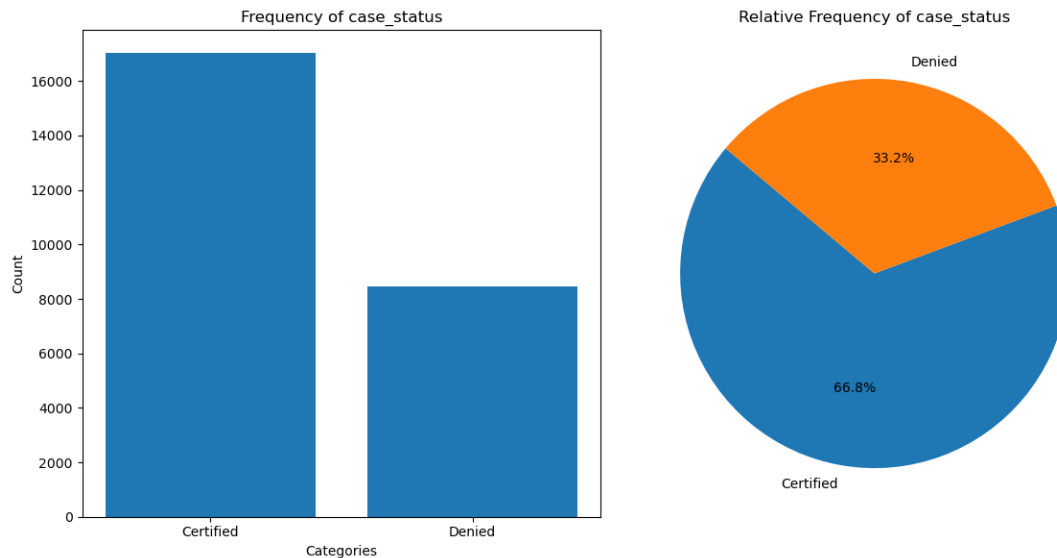


case\_status frequency table

	Class	Frequency
0	Certified	17018
1	Denied	8462

case\_status Relative frequency table

	Class	Frequency	Relative Frequency %
0	Certified	17018	66.79
1	Denied	8462	33.21



## 10 6. Handling Missing Values in Numerical Columns

```
[8]: visa_df.isnull().sum()
```

```
[8]: continent          0
     education_of_employee  0
     has_job_experience    0
     requires_job_training  0
     no_of_employees       0
     yr_of_estab           0
     region_of_employment  0
     prevailing_wage       0
     unit_of_wage          0
     full_time_position    0
     case_status           0
     dtype: int64
```

```
[9]: visa_df.describe()
```

```
[9]:
```

	no_of_employees	yr_of_estab	prevailing_wage
count	25480.000000	25480.000000	25480.000000
mean	5667.043210	1979.409929	74455.814592
std	22877.928848	42.366929	52815.942327
min	-26.000000	1800.000000	2.136700
25%	1022.000000	1976.000000	34015.480000
50%	2109.000000	1997.000000	70308.210000
75%	3504.000000	2005.000000	107735.512500

max            602069.000000    2016.000000    319210.270000

## 11    7. Numerical - UniVariate - Analysis - Using -Pipeline

```
[10]: class HistBoxChartTransformer(BaseEstimator, TransformerMixin):
        def __init__(self):
            pass

        def fit(self, X, y=None):
            return self

        def transform(self, X):
            df=X.copy()
            # getting num cols
            num_cols = df.select_dtypes(exclude='object').columns
            for con_var in num_cols:

                # Create a figure and axes object
                fig, axes = plt.subplots(1, 2, figsize=(14, 6))

                # Plot histogram without KDE on the left
                axes[0].hist(df[con_var], color='skyblue', edgecolor='black')
                axes[0].set_xlabel('Value')
                axes[0].set_ylabel('Frequency')
                axes[0].set_title(f'Histogram {con_var}')


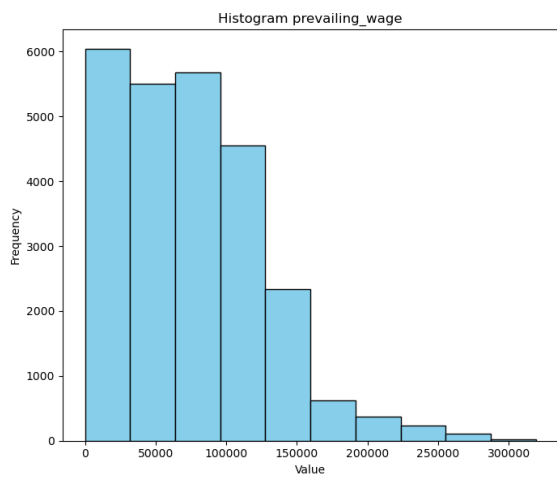
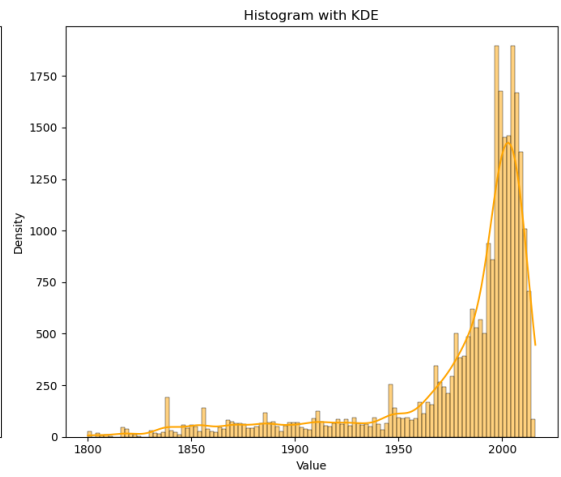
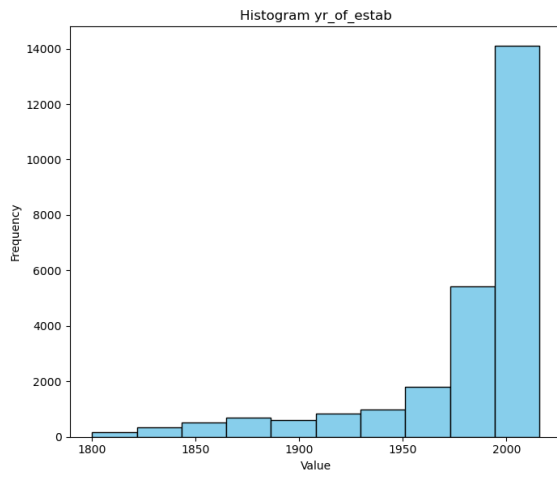
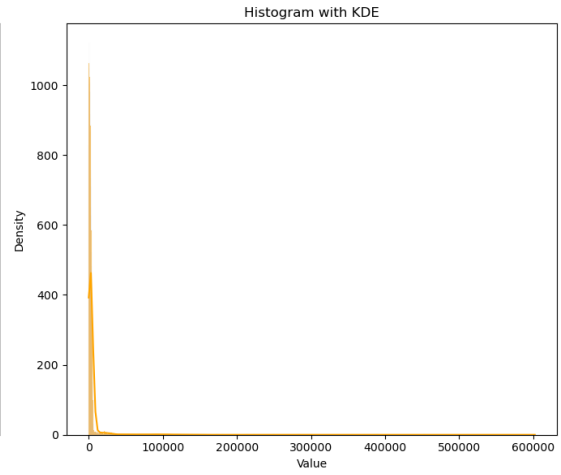
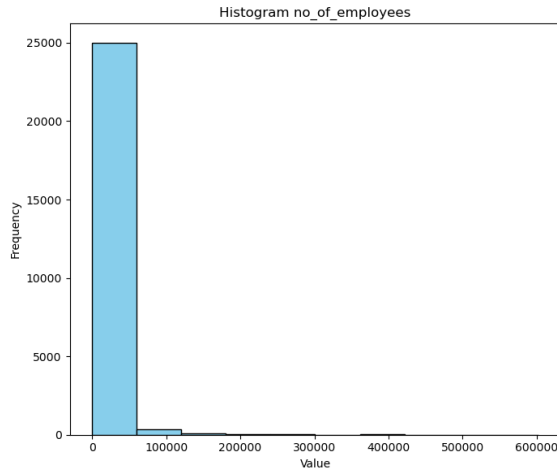
                # Plot histogram with KDE on the right
                sns.histplot(data=df, x=con_var, kde=True, color='orange',
                ↪edgecolor='black', ax=axes[1])
                axes[1].set_xlabel('Value')
                axes[1].set_ylabel('Density')
                axes[1].set_title('Histogram with KDE')

                # Adjust layout
                plt.tight_layout()

                # Show the combined plot
                plt.show()
```

```
[11]: pipeline_num_var = Pipeline([
        ('num_uni_variate_analysis', HistBoxChartTransformer())
    ])

    # Fit and transform your data using the pipeline
    processed_data = pipeline_num_var.fit_transform(visa_df)
```



## 12 8. Numerical - Variables -Outliers Analysis - fillit

## 13 9. Bi Variate Analysis

### 13.1 9.1 cat to target(cat)

```
[12]: cat_vars = visa_df.select_dtypes(include="object").columns
      print(cat_vars)
```

```
Index(['continent', 'education_of_employee', 'has_job_experience',
      'requires_job_training', 'region_of_employment', 'unit_of_wage',
      'full_time_position', 'case_status'],
      dtype='object')
```

```
[13]: target="case_status"
      fig,ax = plt.subplots(4,2,figsize=(15,15))
      for axi,x in zip(ax.flat,cat_vars):
          col1=visa_df[x]
          col2=visa_df[target]
          result=pd.crosstab(col1,col2)
          print(result)
          print("=====")
          result.plot(kind='bar',ax=axi)
```

case_status	Certified	Denied
continent		
Africa	397	154
Asia	11012	5849
Europe	2957	775
North America	2037	1255
Oceania	122	70
South America	493	359

=====

case_status	Certified	Denied
education_of_employee		
Bachelor's	6367	3867
Doctorate	1912	280
High School	1164	2256
Master's	7575	2059

=====

case_status	Certified	Denied
has_job_experience		
N	5994	4684
Y	11024	3778

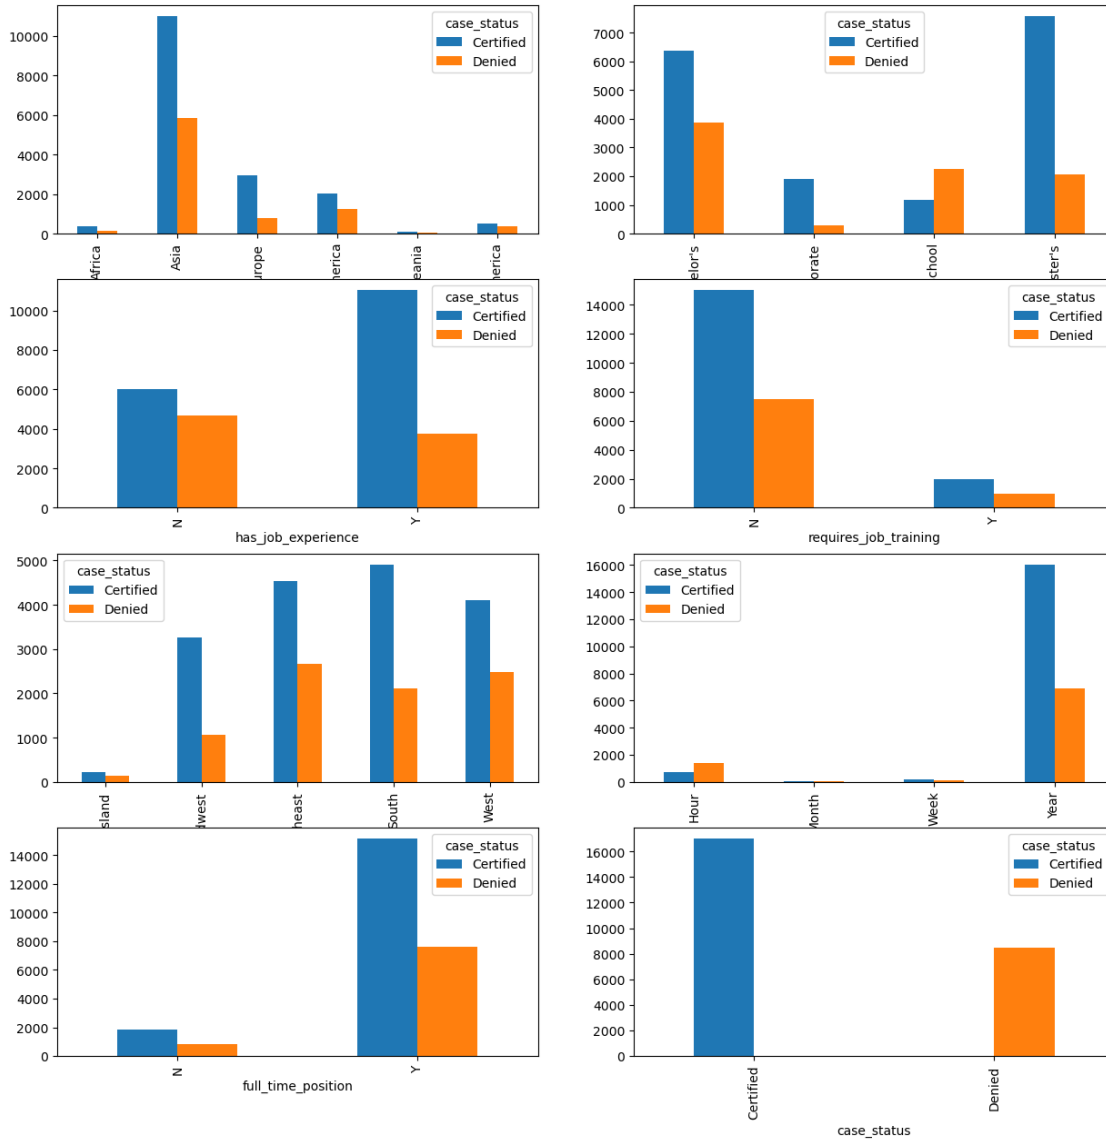
=====

case_status	Certified	Denied
-------------	-----------	--------

```

requires_job_training
N                15012    7513
Y                 2006     949
=====
case_status      Certified Denied
region_of_employment
Island           226     149
Midwest          3253    1054
Northeast        4526    2669
South            4913    2104
West             4100    2486
=====
case_status      Certified Denied
unit_of_wage
Hour             747     1410
Month            55       34
Week             169     103
Year            16047    6915
=====
case_status      Certified Denied
full_time_position
N                1855     852
Y               15163    7610
=====
case_status      Certified Denied
case_status
Certified        17018      0
Denied           0      8462
=====

```




## 13.2 9.2 Num vs Num

```
[14]: num_cols1 = visa_df.select_dtypes(exclude="object").columns.to_list()
      num_cols2 = num_cols1.copy()
      num_cols2
```


```
[14]: ['no_of_employees', 'yr_of_estab', 'prevailing_wage']
```


```
[15]: for i in num_cols1:
      for j in num_cols2:
          if i == j:
              pass
```

```
else:  
    sns.scatterplot(x=i,y=j,hue=target,data=visa_df)  
    plt.show()  
num_cols2.pop(0)
```

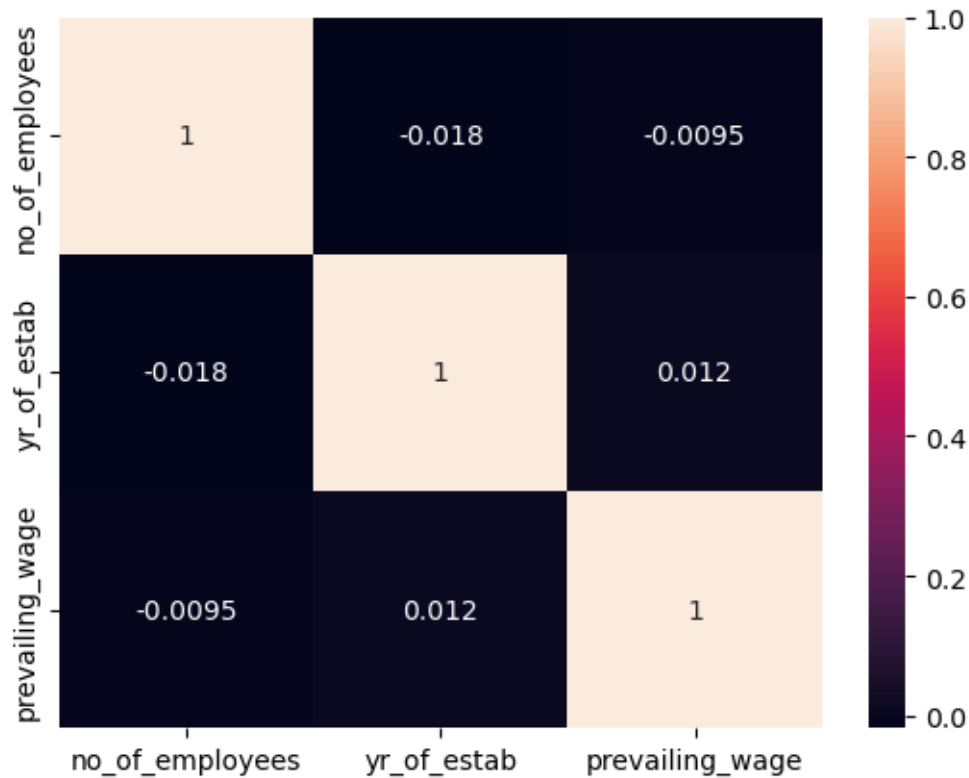






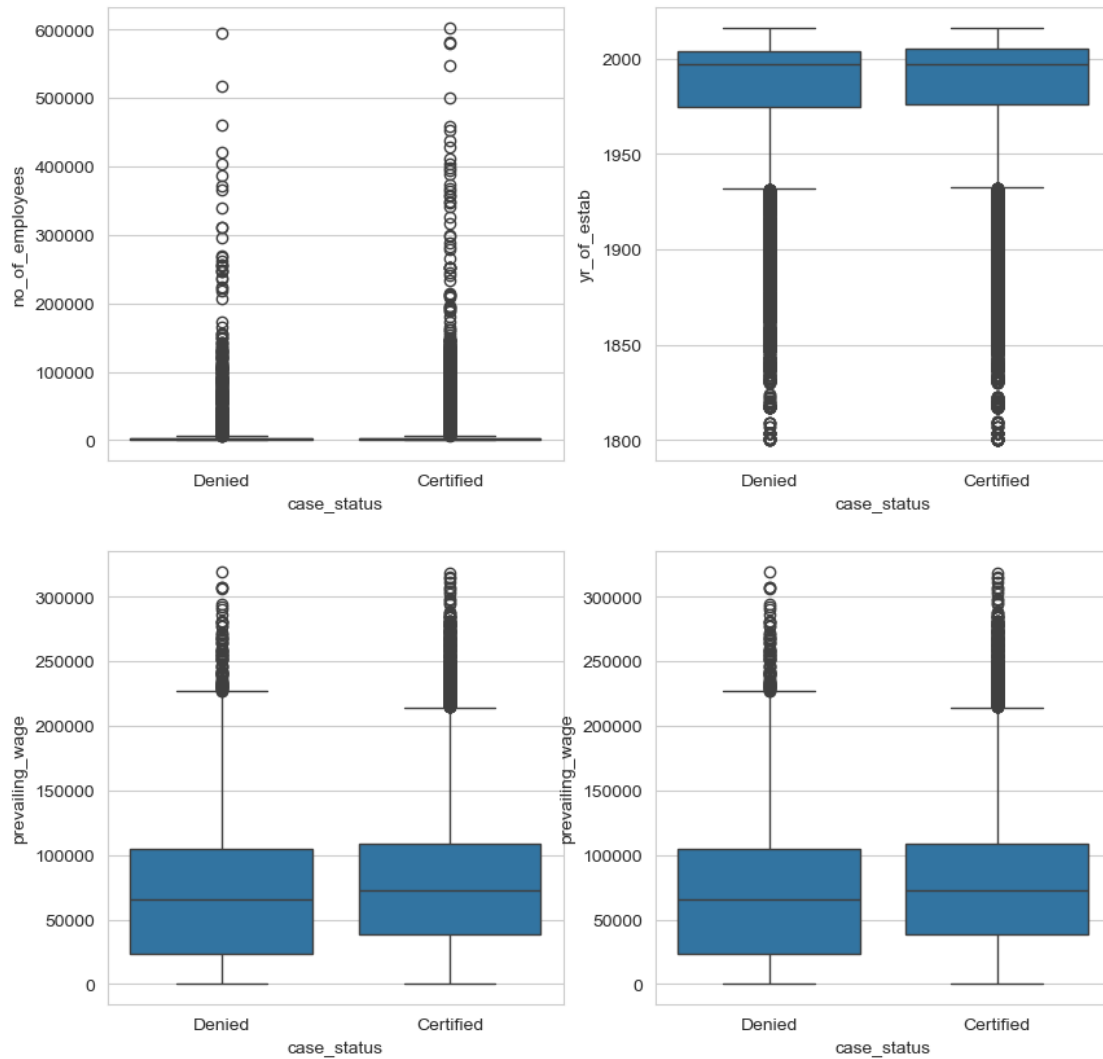


```
[16]: corr_mat=visa_df.corr(numeric_only=True)
sns.heatmap(corr_mat,annot=True)
plt.show()
```



```
[17]: output_var=target
sns.set_style("whitegrid")
fig, axes = plt.subplots(2, 2, figsize=(10, 10))
fig.suptitle('Box-Plots Features Vs Visa Status')
sns.boxplot(ax=axes[0, 0], x=output_var, y='no_of_employees', data=visa_df)
sns.boxplot(ax=axes[0, 1], x=output_var, y='yr_of_estab', data=visa_df)
sns.boxplot(ax=axes[1, 0], x=output_var, y='prevailing_wage', data=visa_df)
sns.boxplot(ax=axes[1, 1], x=output_var, y='prevailing_wage', data=visa_df)
plt.show()
```

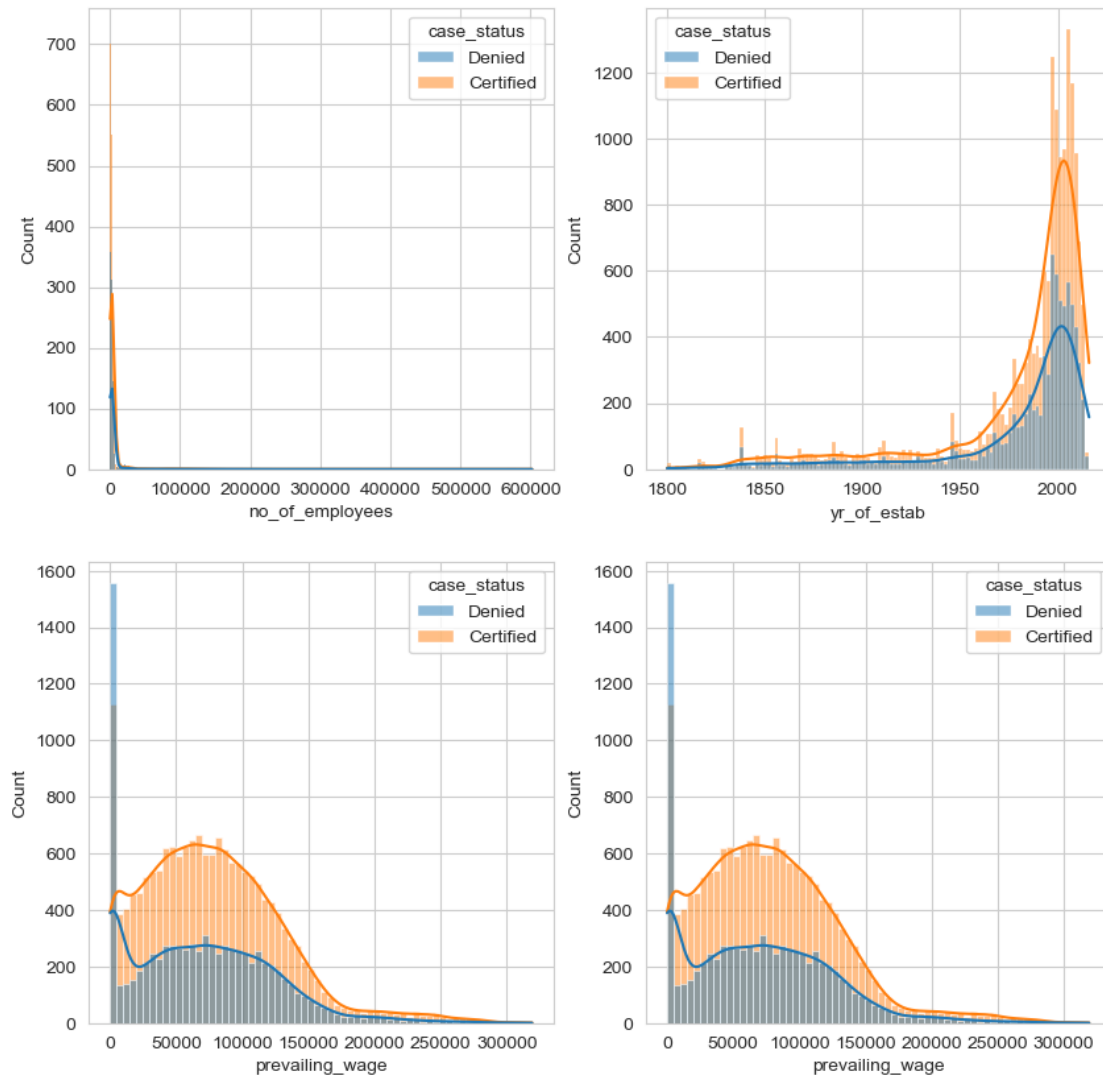
Box-Plots Features Vs Visa Status



```
[18]: sns.set_style("whitegrid")
fig, axes = plt.subplots(2, 2, figsize=(10, 10))
fig.suptitle('Kde-Plots')
sns.histplot(ax=axes[0, 0], hue=output_var, x='no_of_employees',
             data=visa_df, kde=True)
sns.histplot(ax=axes[0, 1], hue=output_var, x='yr_of_estab',
             data=visa_df, kde=True)
sns.histplot(ax=axes[1, 0], hue=output_var, x='prevailing_wage',
             data=visa_df, kde=True)
```

```
sns.histplot(ax=axes[1, 1], hue=output_var, x='prevailing_wage',  
             data=visa_df, kde=True)  
plt.show()
```

Kde-Plots



## 14 10. Data Transformation

```
[19]: visa_df.select_dtypes(exclude='object').columns
```

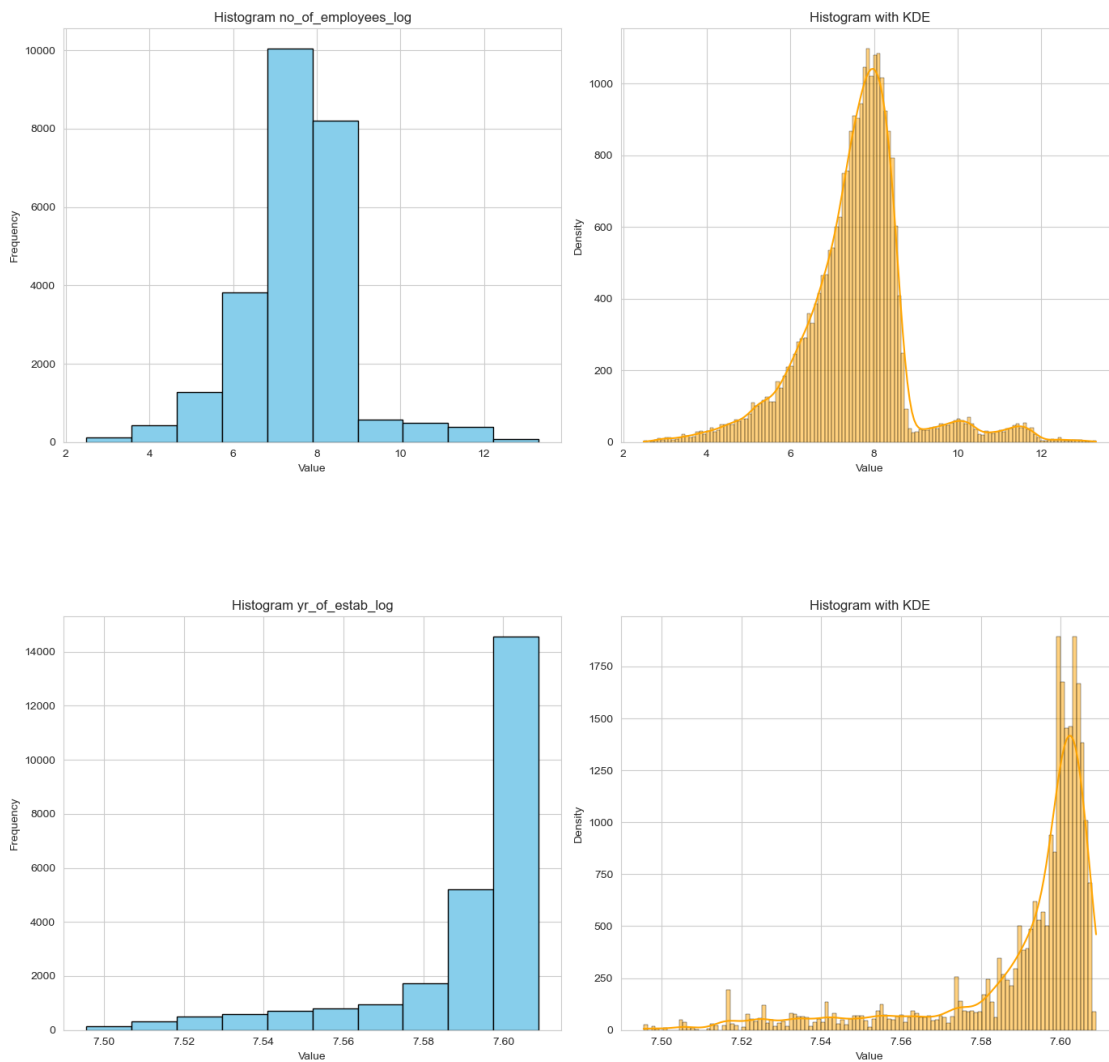
```
[19]: Index(['no_of_employees', 'yr_of_estab', 'prevailing_wage'], dtype='object')
```

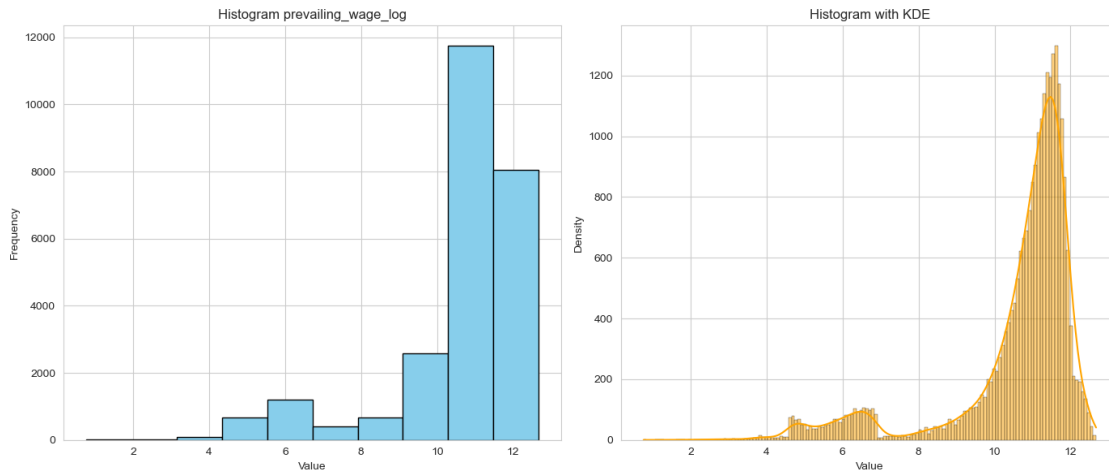
```
[20]: visa_df["no_of_employees_log"]=np.log(visa_df["no_of_employees"])
visa_df["yr_of_estab_log"]=np.log(visa_df["yr_of_estab"])
visa_df["prevailing_wage_log"]=np.log(visa_df["prevailing_wage"])

visa_num_df = visa_df[['no_of_employees_log','yr_of_estab_log',
↳ 'prevailing_wage_log']].copy()
# Fit and transform your data using the pipeline
processed_data = pipeline_num_var.fit_transform(visa_num_df)
```

C:\Users\srishanm\AppData\Local\anaconda3\Lib\site-packages\pandas\core\arraylike.py:396: RuntimeWarning: invalid value encountered in log

```
result = getattr(ufunc, method)(*inputs, **kwargs)
```



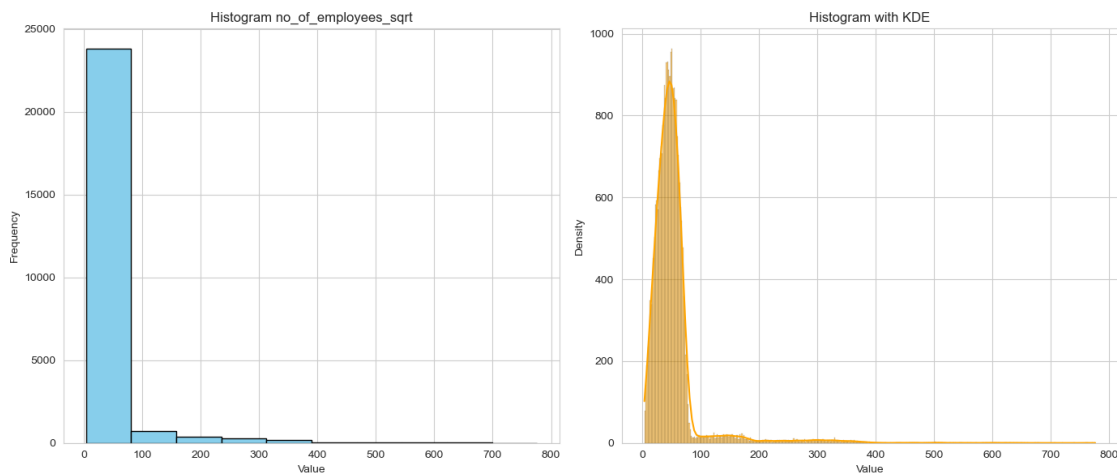


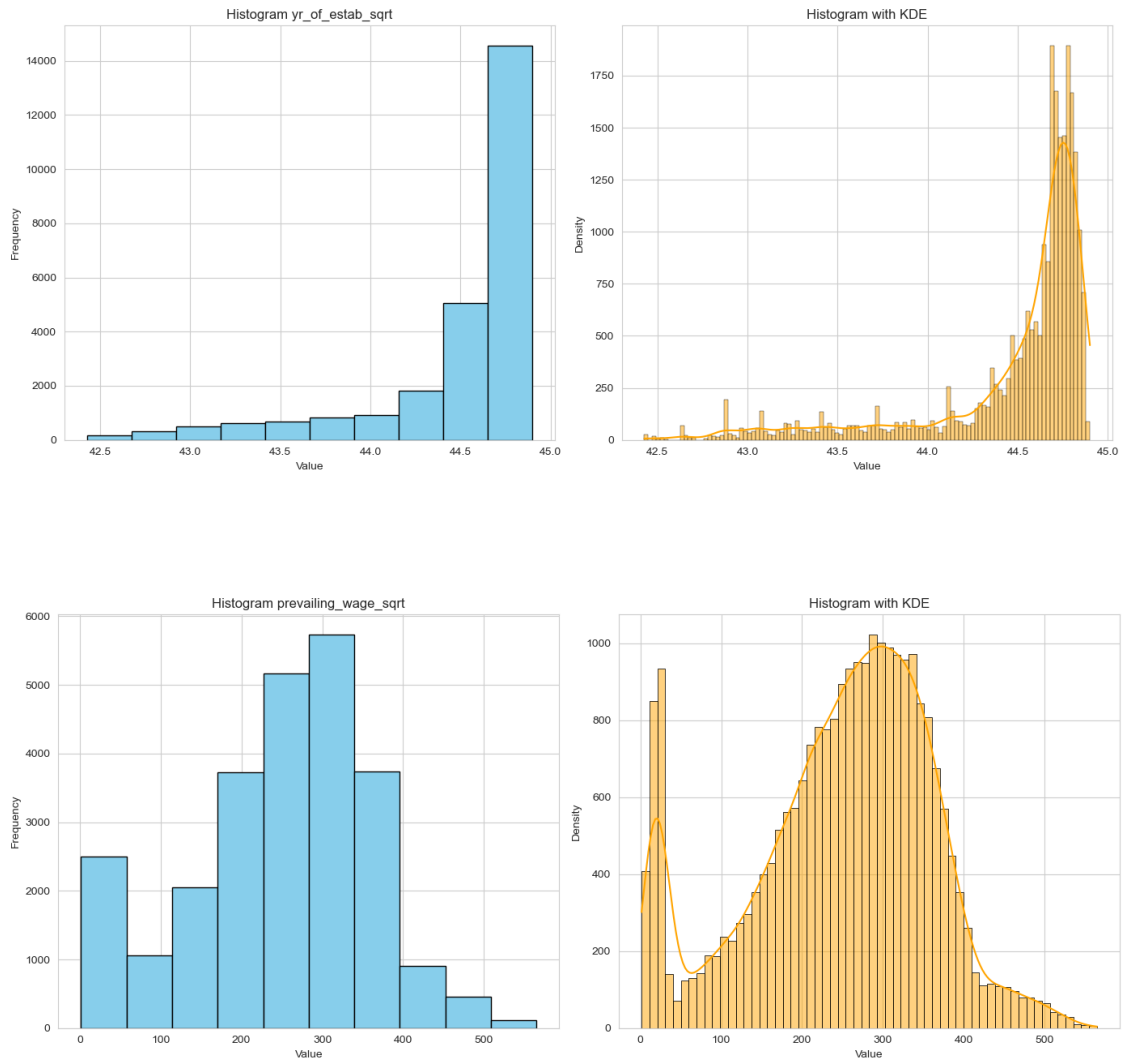
```
[21]: visa_df["no_of_employees_sqrt"]=np.sqrt(visa_df["no_of_employees"])
visa_df["yr_of_estab_sqrt"]=np.sqrt(visa_df["yr_of_estab"])
visa_df["prevailing_wage_sqrt"]=np.sqrt(visa_df["prevailing_wage"])
```

```
visa_num_df = visa_df[['no_of_employees_sqrt','yr_of_estab_sqrt',
↳ 'prevailing_wage_sqrt']].copy()
# Fit and transform your data using the pipeline
processed_data = pipeline_num_var.fit_transform(visa_num_df)
```

C:\Users\srishanm\AppData\Local\anaconda3\Lib\site-packages\pandas\core\arraylike.py:396: RuntimeWarning: invalid value encountered in sqrt

```
result = getattr(ufunc, method)(*inputs, **kwargs)
```

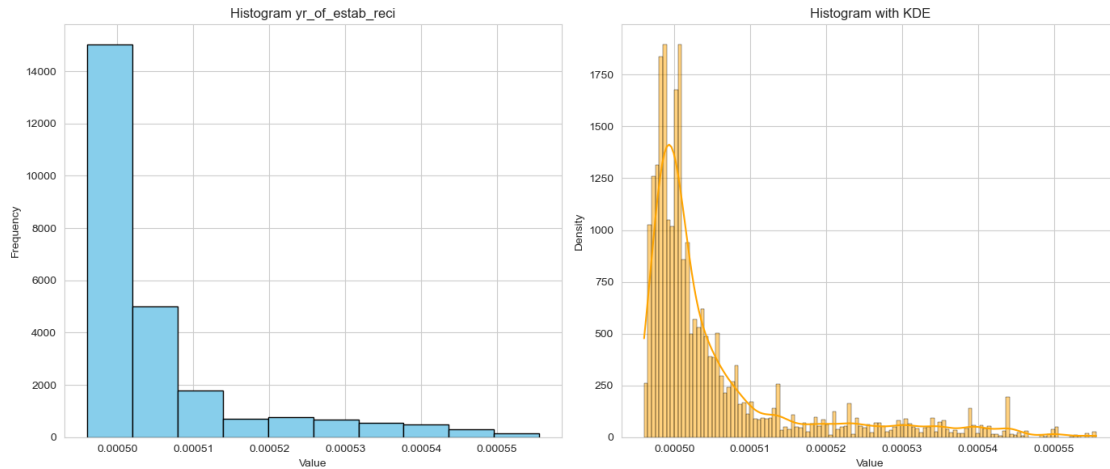




```
[22]: visa_df["yr_of_estab_recipro"] = 1 / (visa_df["yr_of_estab"])
```

```
[23]: visa_num_df = visa_df[['yr_of_estab_recipro']].copy()
      # Fit and transform your data using the pipeline
      processed_data = pipeline_num_var.fit_transform(visa_num_df)
```





## 15 11. Standization - Normalization

```
[24]: scaler = StandardScaler()

# Fit and transform the scaler on the selected columns
scaled_columns = scaler.
    ↪ fit_transform(visa_df[['no_of_employees_log', 'yr_of_estab_log',
    ↪ 'prevailing_wage_sqrt']])

# Replace the original columns with the scaled columns
visa_df[['no_of_employees_log_stand', 'yr_of_estab_log_stand',
    ↪ 'prevailing_wage_sqrt_stand']] = scaled_columns

print(visa_df)
```

	continent	education_of_employee	has_job_experience	\
0	Asia	High School	N	
1	Asia	Master's	Y	
2	Asia	Bachelor's	N	
3	Asia	Bachelor's	N	
4	Africa	Master's	Y	
...	...	...	...	
25475	Asia	Bachelor's	Y	
25476	Asia	High School	Y	
25477	Asia	Master's	Y	
25478	Asia	Master's	Y	
25479	Asia	Bachelor's	Y	

	requires_job_training	no_of_employees	yr_of_estab	\
0	N	14513	2007	

1	N	2412	2002
2	Y	44444	2008
3	N	98	1897
4	N	1082	2005
...	...	...	...
25475	Y	2601	2008
25476	N	3274	2006
25477	N	1121	1910
25478	Y	1918	1887
25479	N	3195	1960

	region_of_employment	prevailing_wage	unit_of_wage	full_time_position	\
0	West	592.2029	Hour		Y
1	Northeast	83425.6500	Year		Y
2	West	122996.8600	Year		Y
3	West	83434.0300	Year		Y
4	South	149907.3900	Year		Y
...	...	...	...	...	
25475	South	77092.5700	Year		Y
25476	Northeast	279174.7900	Year		Y
25477	South	146298.8500	Year		N
25478	West	86154.7700	Year		Y
25479	Midwest	70876.9100	Year		Y

	...	no_of_employees_log	yr_of_estab_log	prevailing_wage_log	\
0	...	9.582800	7.604396	6.383849	
1	...	7.788212	7.601902	11.331711	
2	...	10.701985	7.604894	11.719914	
3	...	4.584967	7.548029	11.331812	
4	...	6.986566	7.603399	11.917773	
...	...	...	...	...	
25475	...	7.863651	7.604894	11.252762	
25476	...	8.093768	7.603898	12.539593	
25477	...	7.021976	7.554859	11.893407	
25478	...	7.559038	7.542744	11.363901	
25479	...	8.069342	7.580700	11.168700	

	no_of_employees_sqrt	yr_of_estab_sqrt	prevailing_wage_sqrt	\
0	120.469913	44.799554	24.335219	
1	49.112117	44.743715	288.834987	
2	210.817457	44.810713	350.709082	
3	9.899495	43.554563	288.849494	
4	32.893768	44.777226	387.178757	
...	...	...	...	
25475	51.000000	44.810713	277.655488	
25476	57.218878	44.788391	528.369937	
25477	33.481338	43.703547	382.490327	
25478	43.794977	43.439613	293.521328	

25479	56.524331	44.271887	266.227177
-------	-----------	-----------	------------

	yr_of_estab_reci	no_of_employees_log_stand	yr_of_estab_log_stand \
0	0.000498	1.633418	0.643017
1	0.000500	0.200370	0.529088
2	0.000498	2.527131	0.665769
3	0.000527	-2.357543	-1.931516
4	0.000499	-0.439774	0.597480
...	...	...	...
25475	0.000498	0.260612	0.665769
25476	0.000499	0.444369	0.620254
25477	0.000524	-0.411498	-1.619582
25478	0.000530	0.017367	-2.172923
25479	0.000510	0.424864	-0.439305

	prevailing_wage_sqrt_stand
0	-1.990753
1	0.357791
2	0.907183
3	0.357920
4	1.231004
...	...
25475	0.258526
25476	2.484668
25477	1.189375
25478	0.399402
25479	0.157052

[25480 rows x 21 columns]

## 16 12. Convert Cat - to - Numerical Columns

```
[25]: cat_vars = visa_df.select_dtypes(include='object').columns
cat_vars
```

```
[25]: Index(['continent', 'education_of_employee', 'has_job_experience',
           'requires_job_training', 'region_of_employment', 'unit_of_wage',
           'full_time_position', 'case_status'],
          dtype='object')
```

```
[26]: from sklearn.preprocessing import LabelEncoder
for var in cat_vars:
    le = LabelEncoder()
    visa_df[var]=le.fit_transform(visa_df[var])
```

## 17 13. SMOTE for Balancing Data

```
[27]: visa_df.columns
```

```
[27]: Index(['continent', 'education_of_employee', 'has_job_experience',  
        'requires_job_training', 'no_of_employees', 'yr_of_estab',  
        'region_of_employment', 'prevailing_wage', 'unit_of_wage',  
        'full_time_position', 'case_status', 'no_of_employees_log',  
        'yr_of_estab_log', 'prevailing_wage_log', 'no_of_employees_sqrt',  
        'yr_of_estab_sqrt', 'prevailing_wage_sqrt', 'yr_of_estab_reci',  
        'no_of_employees_log_stand', 'yr_of_estab_log_stand',  
        'prevailing_wage_sqrt_stand'],  
        dtype='object')
```

```
[28]: visa_df.dropna(inplace=True)
```

```
[29]: Y=visa_df["case_status"]  
X=visa_df[['continent', 'education_of_employee', 'has_job_experience',  
        'requires_job_training',  
        'region_of_employment', 'unit_of_wage',  
        'full_time_position', 'no_of_employees_log_stand',  
        'yr_of_estab_log_stand',  
        'prevailing_wage_sqrt_stand']]  
print(len(Y),len(X))  
print(X.columns)
```

```
25447 25447
```

```
Index(['continent', 'education_of_employee', 'has_job_experience',  
        'requires_job_training', 'region_of_employment', 'unit_of_wage',  
        'full_time_position', 'no_of_employees_log_stand',  
        'yr_of_estab_log_stand', 'prevailing_wage_sqrt_stand'],  
        dtype='object')
```

```
[30]: X, Y = SMOTE().fit_resample(X, Y)  
print(X.columns)  
print(len(Y),len(X))
```

```
Index(['continent', 'education_of_employee', 'has_job_experience',  
        'requires_job_training', 'region_of_employment', 'unit_of_wage',  
        'full_time_position', 'no_of_employees_log_stand',  
        'yr_of_estab_log_stand', 'prevailing_wage_sqrt_stand'],  
        dtype='object')
```

```
34002 34002
```

```
[31]: Y.value_counts()
```

```
[31]: case_status  
1      17001
```

```
0    17001
Name: count, dtype: int64
```

## 18 ML Models

```
[32]: X_train, X_test, Y_train, Y_test = train_test_split(X,Y,test_size=0.
      ↪2,random_state = 42)
```

```
[33]: def draw_heatmap(conf_matrix):
      sns.heatmap(conf_matrix, annot=True)
      plt.xlabel('Predicted Labels')
      plt.ylabel('Actual Labels')
      plt.title('Confusion Matrix')
      plt.show()
```

### 18.1 Logistic Regression

```
[34]: lg_model = LogisticRegression(solver='saga', max_iter=500, random_state=42)
      lg_model.fit(X_train, Y_train)

      print("Model - Logistic Regression")
      score = lg_model.score(X_train, Y_train)
      print('accuracy train score overall :', score)
      score = lg_model.score(X_test, Y_test)
      print('accuracy test score overall :', score)

      y_pred = lg_model.predict(X_test)
      print(classification_report(Y_test, y_pred))
      print(confusion_matrix(Y_test, y_pred))
      conf_matrix = confusion_matrix(Y_test, y_pred)
      draw_heatmap(conf_matrix)
```


Model - Logistic Regression

accuracy train score overall : 0.632182640344105

accuracy test score overall : 0.6228495809439788

	precision	recall	f1-score	support
0	0.60	0.68	0.64	3364
1	0.65	0.56	0.60	3437
accuracy			0.62	6801
macro avg	0.63	0.62	0.62	6801
weighted avg	0.63	0.62	0.62	6801

```
[[2302 1062]
 [1503 1934]]
```



## 18.2 GaussianNB

```
[35]: from sklearn.naive_bayes import GaussianNB, CategoricalNB
gnb_model = GaussianNB()
gnb_model.fit(X_train,Y_train)

print("Model-GaussianNB")
print("train score",gnb_model.score(X_train,Y_train))
print("test score",gnb_model.score(X_test,Y_test))

y_pred = gnb_model.predict(X_test)
print(classification_report(Y_test, y_pred))
print(confusion_matrix(Y_test, y_pred))
conf_matrix = confusion_matrix(Y_test, y_pred)
draw_heatmap(conf_matrix)
```

Model-GaussianNB

train score 0.5835447226204918


test score 0.5744743420085282

	precision	recall	f1-score	support
--	-----------	--------	----------	---------

0	0.54	0.87	0.67	3364
1	0.69	0.29	0.41	3437
accuracy			0.57	6801
macro avg	0.62	0.58	0.54	6801
weighted avg	0.62	0.57	0.54	6801

```
[[2912  452]
 [2442  995]]
```



## 19 Support Vector Machine - Classifier

```
[36]: from sklearn.svm import SVC
      # Initialize the SVM classifier
      svm_linear_classifier = SVC(kernel='linear', random_state=42)

      # Train the SVM classifier
      svm_linear_classifier.fit(X_train, Y_train)
      print("model-Support Vector Machine - kernel - linear -Classifier")
```

```

y_pred = svm_linear_classifier.predict(X_train)
# Calculate the accuracy of the model
accuracy = accuracy_score(Y_train, y_pred)
print("Train Accuracy:", accuracy)

# Predict the classes for test set
y_pred = svm_linear_classifier.predict(X_test)
# Calculate the accuracy of the model
accuracy = accuracy_score(Y_test, y_pred)
print("Test Accuracy:", accuracy)

print(classification_report(Y_test, y_pred))
print(confusion_matrix(Y_test, y_pred))
conf_matrix = confusion_matrix(Y_test, y_pred)
draw_heatmap(conf_matrix)

```

model-Suport Vector Machine - kernel - linear -Classifier

Train Accuracy: 0.6231388551891475

Test Accuracy: 0.6194677253345097


	precision	recall	f1-score	support
0	0.61	0.62	0.62	3364
1	0.63	0.62	0.62	3437
accuracy			0.62	6801
macro avg	0.62	0.62	0.62	6801
weighted avg	0.62	0.62	0.62	6801

```

[[2094 1270]
 [1318 2119]]

```





```
[37]: svm_rbf_classifier = SVC(kernel='rbf', random_state=42)

# Train the SVM classifier
svm_rbf_classifier.fit(X_train, Y_train)
print("model-Suport Vector Machine - Kernel -rbf - Classifier")
y_pred = svm_rbf_classifier.predict(X_train)
# Calculate the accuracy of the model
accuracy = accuracy_score(Y_train, y_pred)
print("Train Accuracy:", accuracy)


# Predict the classes for test set
y_pred = svm_rbf_classifier.predict(X_test)
# Calculate the accuracy of the model
accuracy = accuracy_score(Y_test, y_pred)
print("Test Accuracy:", accuracy)

print(classification_report(Y_test, y_pred))
print(confusion_matrix(Y_test, y_pred))
conf_matrix = confusion_matrix(Y_test, y_pred)
draw_heatmap(conf_matrix)
```

model-Suport Vector Machine - Kernel -rbf - Classifier  
 Train Accuracy: 0.6752692915701629  
 Test Accuracy: 0.6800470519041317

	precision	recall	f1-score	support
0	0.67	0.70	0.68	3364
1	0.69	0.67	0.68	3437
accuracy			0.68	6801
macro avg	0.68	0.68	0.68	6801
weighted avg	0.68	0.68	0.68	6801

```
[[2338 1026]
 [1150 2287]]
```



```
[38]: svm_poly_classifier = SVC(kernel='poly', random_state=42)

# Train the SVM classifier
svm_poly_classifier.fit(X_train, Y_train)
print("model-Suport Vector Machine - Kernel -poly - Classifier")
```

```

y_pred = svm_poly_classifier.predict(X_train)
# Calculate the accuracy of the model
accuracy = accuracy_score(Y_train, y_pred)
print("Train Accuracy:", accuracy)

# Predict the classes for test set
y_pred = svm_poly_classifier.predict(X_test)

# Calculate the accuracy of the model
accuracy = accuracy_score(Y_test, y_pred)
print("Test Accuracy:", accuracy)

print(classification_report(Y_test, y_pred))
print(confusion_matrix(Y_test, y_pred))
conf_matrix = confusion_matrix(Y_test, y_pred)
draw_heatmap(conf_matrix)

```

model-Suport Vector Machine - Kernel -poly - Classifier

Train Accuracy: 0.6733575971471637


Test Accuracy: 0.6796059403028967

	precision	recall	f1-score	support
0	0.67	0.70	0.68	3364
1	0.69	0.66	0.67	3437
accuracy			0.68	6801
macro avg	0.68	0.68	0.68	6801
weighted avg	0.68	0.68	0.68	6801

```

[[2361 1003]
 [1176 2261]]

```



## 19.1 Decision Tree

```
[39]: dt_clf = DecisionTreeClassifier(max_leaf_nodes=20,random_state=42)
dt_clf.fit(X_train, Y_train)
print("Model-Decion Tree")

accuracy=dt_clf.score(X_train, Y_train)
print(f"train score: {accuracy}")


accuracy=dt_clf.score(X_test, Y_test)
print(f"test score: {accuracy}")

y_pred=dt_clf.predict(X_test)
print(classification_report(Y_test, y_pred))
print(confusion_matrix(Y_test, y_pred))
conf_matrix = confusion_matrix(Y_test, y_pred)
draw_heatmap(conf_matrix)
```

```
Model-Decion Tree
train score: 0.7029888607036506
test score: 0.709748566387296
```

	precision	recall	f1-score	support
0	0.69	0.75	0.72	3364
1	0.73	0.67	0.70	3437
accuracy			0.71	6801
macro avg	0.71	0.71	0.71	6801
weighted avg	0.71	0.71	0.71	6801

```
[[2508 856]
 [1118 2319]]
```



## 19.2 Random Forest

```
[40]: rf_clf= RandomForestClassifier(n_estimators = 1000, random_state = 42,
    ↪max_leaf_nodes=20)
rf_clf.fit(X_train, Y_train)
print("Model- Random Forest Tree")

accuracy=rf_clf.score(X_train, Y_train)
```

```

print(f"train score: {accuracy}")

accuracy=rf_clf.score(X_test, Y_test)
print(f"test score: {accuracy}")

y_pred=rf_clf.predict(X_test)
print(classification_report(Y_test, y_pred))
print(confusion_matrix(Y_test, y_pred))
conf_matrix = confusion_matrix(Y_test, y_pred)
draw_heatmap(conf_matrix)

```

Model- Random Forest Tree

train score: 0.7099003713098783


test score: 0.7209233936185855

	precision	recall	f1-score	support
0	0.71	0.74	0.73	3364
1	0.74	0.70	0.72	3437
accuracy			0.72	6801
macro avg	0.72	0.72	0.72	6801
weighted avg	0.72	0.72	0.72	6801

```

[[2502  862]
 [1036 2401]]

```



### 19.3 AdaBoost

```
[41]: base_classifier = DecisionTreeClassifier(max_depth=1)
      adaboost_clf = AdaBoostClassifier( n_estimators=50, random_state=42)

      # Train the AdaBoost classifier
      adaboost_clf.fit(X_train, Y_train)

      print("Model-AdaBoost")
      print("train score",adaboost_clf.score(X_train, Y_train))

      # Predict on the test set
      y_pred = adaboost_clf.predict(X_test)

      # Calculate accuracy
      accuracy = accuracy_score(Y_test, y_pred)
      print(f"test score: {accuracy}")

      print(classification_report(Y_test, y_pred))
      print(confusion_matrix(Y_test, y_pred))
```

```
conf_matrix = confusion_matrix(Y_test, y_pred)
draw_heatmap(conf_matrix)
```


Model-AdaBoost

train score 0.703356494246535

test score: 0.7131304219967651

	precision	recall	f1-score	support
0	0.70	0.74	0.72	3364
1	0.73	0.69	0.71	3437
accuracy			0.71	6801
macro avg	0.71	0.71	0.71	6801
weighted avg	0.71	0.71	0.71	6801

```
[[2473  891]
 [1060 2377]]
```





## 19.4 GradientBoostingClassifier

```
[42]: gdb_clf = GradientBoostingClassifier(n_estimators=20, learning_rate=0.1,
    ↪max_depth=1, random_state=42)
gdb_clf.fit(X_train, Y_train)
print("model-Gradient Boosting Classifier")

accuracy = gdb_clf.score(X_train, Y_train)
print("Train Accuracy:", accuracy)

accuracy = gdb_clf.score(X_test, Y_test)
print("Test Accuracy:", accuracy)

print(classification_report(Y_test, y_pred))
print(confusion_matrix(Y_test, y_pred))
conf_matrix = confusion_matrix(Y_test, y_pred)
draw_heatmap(conf_matrix)
```


model-Gradient Boosting Classifier

Train Accuracy: 0.6283592514981067

Test Accuracy: 0.6196147625349213

	precision	recall	f1-score	support
0	0.70	0.74	0.72	3364
1	0.73	0.69	0.71	3437
accuracy			0.71	6801
macro avg	0.71	0.71	0.71	6801
weighted avg	0.71	0.71	0.71	6801

```
[[2473  891]
 [1060 2377]]
```



## 19.5 XGBClassifier

```
[43]: from xgboost import XGBClassifier
xgmodel = XGBClassifier()
xgmodel.fit(X_train, Y_train)

print("model- XGB Classifier")
# Make predictions on the test set
y_pred = xgmodel.predict(X_train)
accuracy = accuracy_score(Y_train, y_pred)
print("Train Accuracy:", accuracy)
# Evaluate the model

# Make predictions on the test set
y_pred = xgmodel.predict(X_test)
accuracy = accuracy_score(Y_test, y_pred)
print("Test Accuracy:", accuracy)

print(classification_report(Y_test, y_pred))
```

```
print(confusion_matrix(Y_test, y_pred))
conf_matrix = confusion_matrix(Y_test, y_pred)
draw_heatmap(conf_matrix)
```


model- XGB Classifier

Train Accuracy: 0.837101577147899

Test Accuracy: 0.7656227025437435

	precision	recall	f1-score	support
0	0.75	0.79	0.77	3364
1	0.78	0.75	0.76	3437
accuracy			0.77	6801
macro avg	0.77	0.77	0.77	6801
weighted avg	0.77	0.77	0.77	6801

```
[[2646  718]
 [ 876 2561]]
```



[ ]: