


Second-Hand-Cars-Price-Prediction

June 1, 2024

- 1 Objective : Price Prediction - On Second Hand Cars
- 2 EDA - Python
- 3 Insights - Patterns
- 4 Regression



5 1. Import Python Modules

```
[1]: # Load necessary python modules
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
from tabulate import tabulate
```

```

from sklearn.base import BaseEstimator, TransformerMixin
from sklearn.pipeline import Pipeline

from sklearn.preprocessing import StandardScaler
from sklearn.preprocessing import MinMaxScaler
from sklearn.model_selection import train_test_split

from sklearn.linear_model import LinearRegression, Lasso, Ridge, ElasticNet

from sklearn.tree import DecisionTreeRegressor
from sklearn.ensemble import RandomForestRegressor
from sklearn.svm import SVR
from sklearn.ensemble import AdaBoostRegressor
from sklearn.ensemble import GradientBoostingRegressor
import xgboost as xgb

from sklearn.metrics import mean_squared_error
from sklearn.metrics import r2_score

```

6 2. Load Cars Price Dataset

```

[2]: file_path = r"Cars-SecondHand.xlsx"
cars_price_df = pd.read_excel(file_path)
cars_price_df

```

```

[2]:
      Brand      Price  Mileage  EngineV  Year
0  Mercedes-Benz  222000.0        1      6.3  2016
1  Mercedes-Benz  177000.0        1      5.5  2016
2  Mercedes-Benz  177777.0        1      5.5  2016
3  Mercedes-Benz  199999.0        1      5.5  2016
4  Mercedes-Benz  199999.0        1      5.5  2016
...      ...      ...      ...      ...
3998      Toyota      600.0       10      1.5  1979
3999  Mercedes-Benz      2990.0     300      2.8  1979
4000  Mercedes-Benz      2300.0     261      2.3  1978
4001  Mercedes-Benz      5500.0     440      2.0  1978
4002  Mercedes-Benz      34999.0    150      2.8  1969

```

[4003 rows x 5 columns]

7 3. Basic Inspection on dataset

```
[3]: def basic_inspection_dataset(table):  
    """Generates a basic inspection dataset from the given table."""  
  
    print("top 5 rows - using head")  
    print(table.head())  
    print()  
  
    print("bottom 5 rows using tail")  
    print(table.tail())  
    print()  
  
    print("numbers of samples and columns")  
    print(table.shape)  
    print()  
  
    print("numbers of samples ")  
    print(len(table))  
    print()  
  
    print("numbers of entries in the data frame")  
    print(table.size)  
    print()  
  
    print("Columns Names")  
    print(table.columns)  
    print()  
  
    print("Columns dtypes")  
    print(table.dtypes)  
    print()  
  
    print("Dataframe info")  
    print(table.info())  
    print()  
  
    print()  
    print("check the missing value in each column")  
    print(table.isnull().sum())  
  
    print()  
    print("check the missing value in each column")  
    print(table.isna().sum())  
  
    print()  
    print("table summary ")
```

```
print(table.describe())
```

```
basic_inspection_dataset(cars_price_df)
```

top 5 rows - using head

	Brand	Price	Mileage	EngineV	Year
0	Mercedes-Benz	222000.0	1	6.3	2016
1	Mercedes-Benz	177000.0	1	5.5	2016
2	Mercedes-Benz	177777.0	1	5.5	2016
3	Mercedes-Benz	199999.0	1	5.5	2016
4	Mercedes-Benz	199999.0	1	5.5	2016

bottom 5 rows using tail

	Brand	Price	Mileage	EngineV	Year
3998	Toyota	600.0	10	1.5	1979
3999	Mercedes-Benz	2990.0	300	2.8	1979
4000	Mercedes-Benz	2300.0	261	2.3	1978
4001	Mercedes-Benz	5500.0	440	2.0	1978
4002	Mercedes-Benz	34999.0	150	2.8	1969

numbers of samples and columns

```
(4003, 5)
```

numbers of samples

```
4003
```

numbers of entries in the data frame

```
20015
```

Columns Names

```
Index(['Brand', 'Price', 'Mileage', 'EngineV', 'Year'], dtype='object')
```

Columns dtypes

```
Brand      object
```

```
Price      float64
```

```
Mileage    int64
```

```
EngineV    float64
```

```
Year       int64
```

```
dtype: object
```

Dataframe info

```
<class 'pandas.core.frame.DataFrame'>
```

```
RangeIndex: 4003 entries, 0 to 4002
```

```
Data columns (total 5 columns):
```

#	Column	Non-Null Count	Dtype
0	Brand	4003 non-null	object
1	Price	4003 non-null	float64

```

2   Mileage  4003 non-null   int64
3   EngineV  4003 non-null   float64
4   Year     4003 non-null   int64
dtypes: float64(2), int64(2), object(1)
memory usage: 156.5+ KB
None

```

check the missing value in each column

```

Brand      0
Price      0
Mileage    0
EngineV    0
Year       0
dtype: int64

```

check the missing value in each column

```

Brand      0
Price      0
Mileage    0
EngineV    0
Year       0
dtype: int64

```

table summary

	Price	Mileage	EngineV	Year
count	4003.000000	4003.000000	4003.000000	4003.000000
mean	19619.014218	163.419935	2.467732	2006.395703
std	25868.124801	103.406160	0.975549	6.695288
min	600.000000	1.000000	0.600000	1969.000000
25%	7000.000000	90.000000	1.800000	2003.000000
50%	11500.000000	158.000000	2.200000	2008.000000
75%	21900.000000	230.000000	3.000000	2012.000000
max	300000.000000	980.000000	6.500000	2016.000000

8 4. Handling Missing Values - Cat

```
[4]: # There is no missing values in cat columns
```

9 5. Categorical- Variable - Analysis -Using Pipeline

```
[2]: class BarPieChartTransformer(BaseEstimator, TransformerMixin):
      def __init__(self):
          pass

      def fit(self, X, y=None):
```

```

        return self

    def transform(self, X):
        df=X.copy()
        cat_cols = df.select_dtypes(include='object').columns
        for cat_name in cat_cols:
            value_counts = df[cat_name].value_counts().reset_index()
            # Rename the columns
            value_counts.columns = ['Class', 'Frequency']

            # Print the result as a table
            print(f"{cat_name} frequency table")
            print(tabulate(value_counts, headers='keys', tablefmt='pretty'))

            # Calculate relative frequency
            total_count = value_counts['Frequency'].sum()
            value_counts['Relative Frequency %'] =
↪round((value_counts['Frequency'] / total_count)*100,2)

            # Print the result as a table
            print(f"{cat_name} Relative frequency table")
            print(tabulate(value_counts, headers='keys', tablefmt='pretty'))

            # Extract the values and index from value counts
            value_counts = df[cat_name].value_counts()
            values = value_counts.values
            labels = value_counts.index

            fig, axs = plt.subplots(1, 2, figsize=(18, 6)) # 1 row, 2 columns
            # Create a bar graph
            axs[0].bar(labels, values)
            axs[0].set_title(f'Frequency of {cat_name}')
            axs[0].set_xlabel('Categories') # Set x-label
            axs[0].set_ylabel('Count') # Set y-label

            axs[1].pie(value_counts.values, labels=value_counts.index,
↪autopct='%0.2f%%', startangle=40)
            axs[1].set_title(f'Relative Frequency of {cat_name}')
            plt.tight_layout()
            # Show the plot
            plt.show()

```

```

[6]: pipeline_cat_var = Pipeline([
    ('bar_pie_chart', BarPieChartTransformer())
])

```

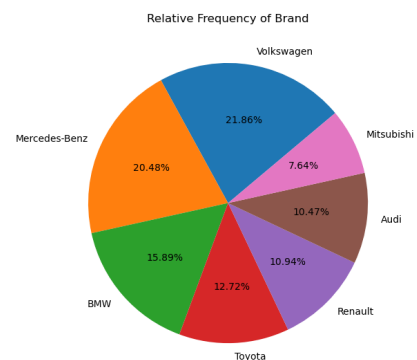
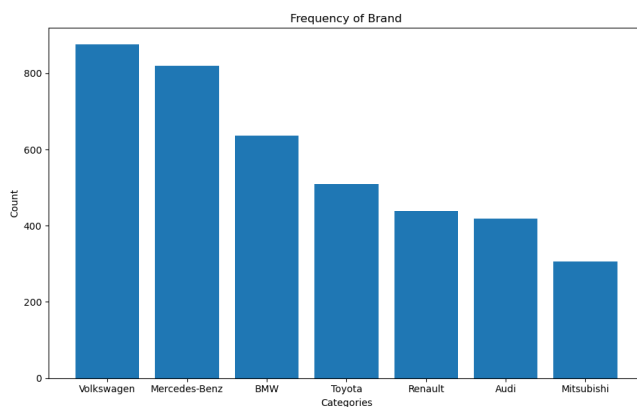
```
# Fit and transform your data using the pipeline
processed_data = pipeline_cat_var.fit_transform(cars_price_df)
```

Brand frequency table

	Class	Frequency
0	Volkswagen	875
1	Mercedes-Benz	820
2	BMW	636
3	Toyota	509
4	Renault	438
5	Audi	419
6	Mitsubishi	306

Brand Relative frequency table

	Class	Frequency	Relative Frequency %
0	Volkswagen	875	21.86
1	Mercedes-Benz	820	20.48
2	BMW	636	15.89
3	Toyota	509	12.72
4	Renault	438	10.94
5	Audi	419	10.47
6	Mitsubishi	306	7.64



Observations 1. Volkswagen with highest numbers of sales with 875 2. Mitsubishi with lowest numbers of sales with 306

10 6. Handling Missing Values in Numerical Columns

```
[7]: # There is no missing values in num columns
```

11 7. Numerical - Variables - Analysis - Using -Pipeline

```
[8]: class HistBoxChartTransformer(BaseEstimator, TransformerMixin):
    def __init__(self):
        pass

    def fit(self, X, y=None):
        return self

    def transform(self, X):
        df=X.copy()
        num_cols = df.select_dtypes(exclude='object').columns
        for con_var in num_cols:

            # Create a figure and axes object
            fig, axes = plt.subplots(1, 2, figsize=(14, 6))

            # Plot histogram without KDE on the left
            axes[0].hist(df[con_var], color='skyblue', edgecolor='black')
            axes[0].set_xlabel('Value')
            axes[0].set_ylabel('Frequency')
            axes[0].set_title(f'Histogram {con_var}')

            # Plot histogram with KDE on the right
            sns.histplot(data=df, x=con_var, kde=True, color='orange',
            ↪edgecolor='black', ax=axes[1])
            axes[1].set_xlabel('Value')
            axes[1].set_ylabel('Density')
            axes[1].set_title('Histogram with KDE')


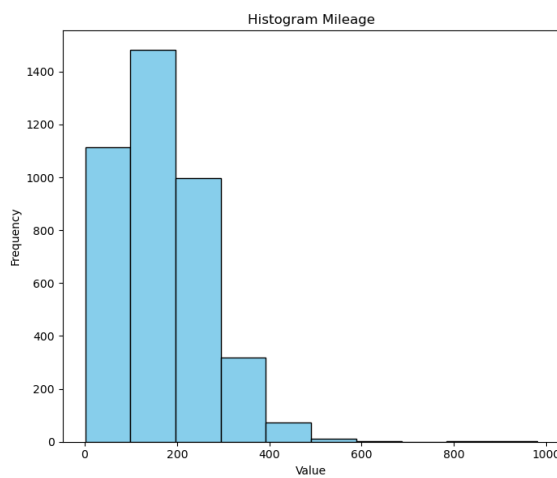
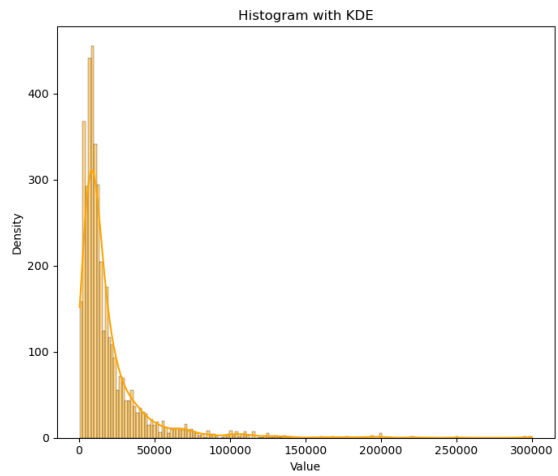
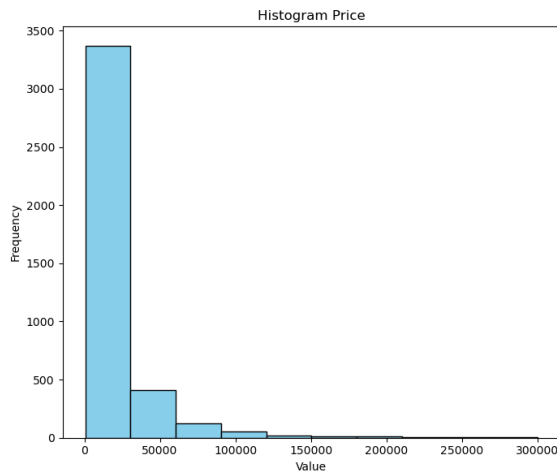
            # Adjust layout
            plt.tight_layout()

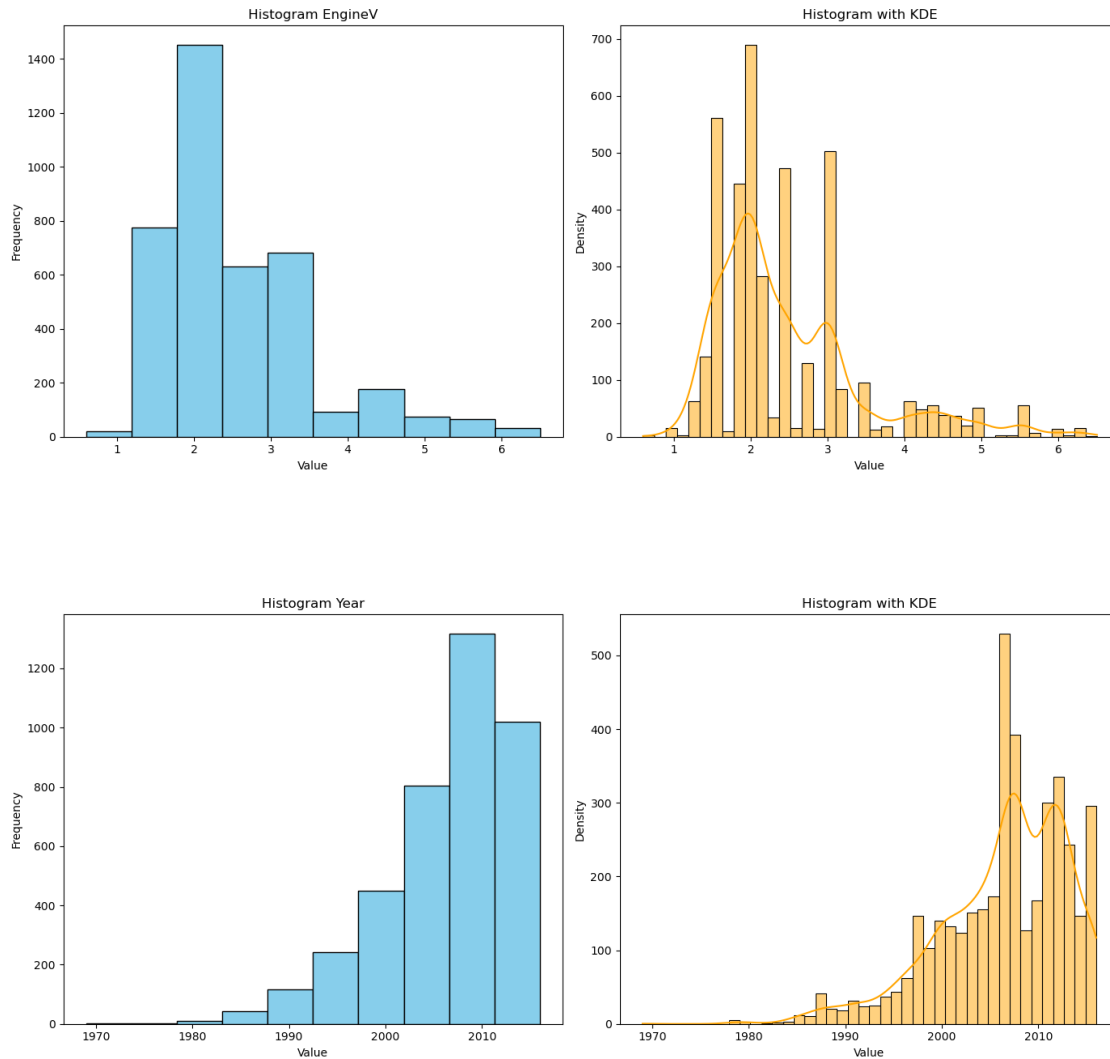
            # Show the combined plot
            plt.show()
```

```
[9]: pipeline_num_var = Pipeline([
    ('hist_box_chart', HistBoxChartTransformer())
])
```



```
cars_price_num_df = cars_price_df[['Price', 'Mileage', 'EngineV', 'Year']]  
# Fit and transform your data using the pipeline  
processed_data = pipeline_num_var.fit_transform(cars_price_num_df)
```





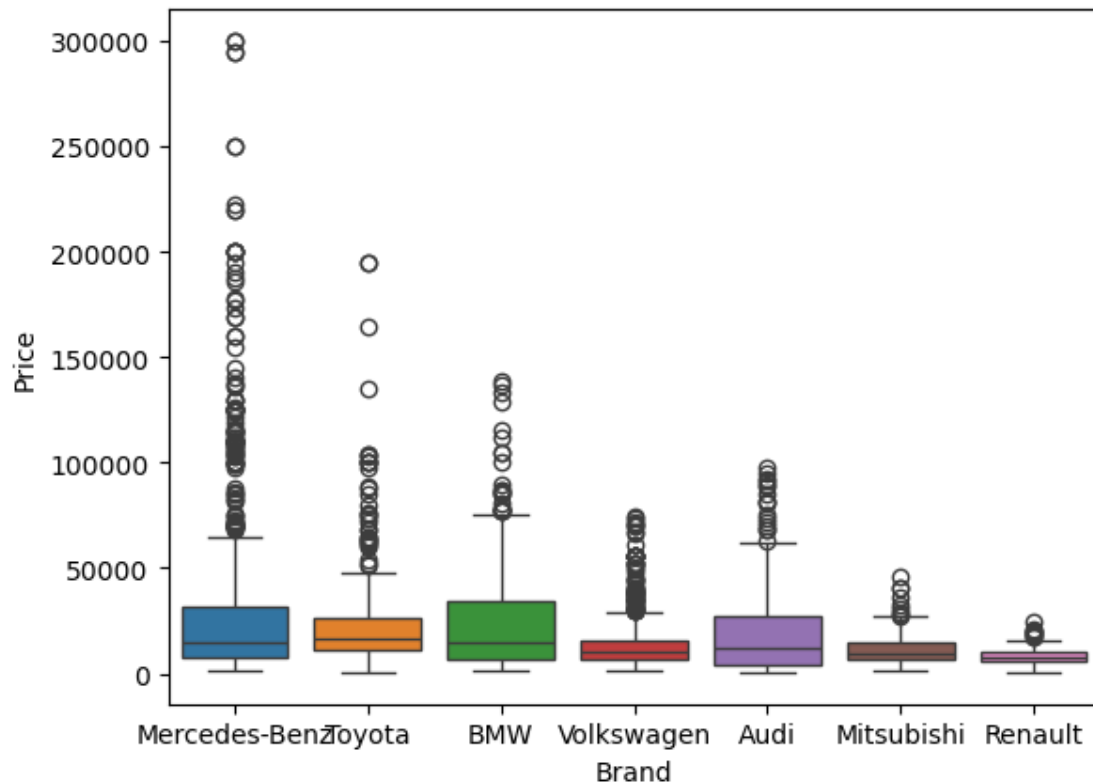
Observations 1. 'Price', 'Mileage', 'EngineV', 'Year' are numerical columns 2. All are not normally distributed

8. Numerical - Variables -Outliers Analysis

12 9. Bi Variate Analysis

12.1 Cat Vs Num


```
[10]: # Create a box plot with hue
sns.boxplot(x='Brand', y='Price', hue='Brand', data=cars_price_df)
# Show the plot
plt.show()
```





Observations 1. Mercedes-Benz with Highest price 2. Renault with lowest price 3. There are 7 brands are available in the data frame in Brand Column

12.2 Num Vs Num

```
[11]: #print(cars_price_df.columns)
for num_var in [ 'Mileage', 'EngineV', 'Year']:
    sns.scatterplot(data=cars_price_df, y='Price', x=num_var)
    plt.show()
```







Observations 1. There is relationship b/w variables to price

```
[12]: cars_price_df[['Price', 'Mileage', 'EngineV', 'Year']].corr()
```

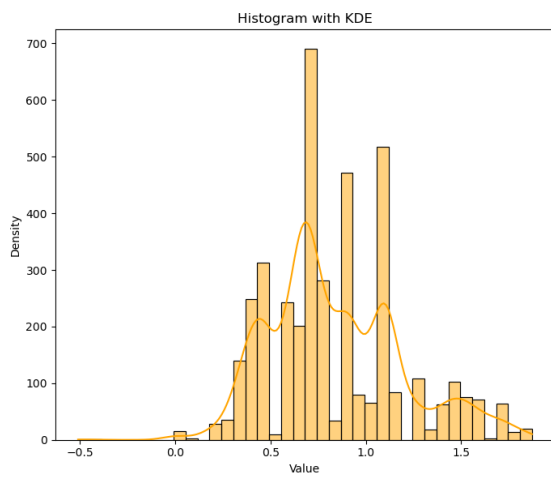
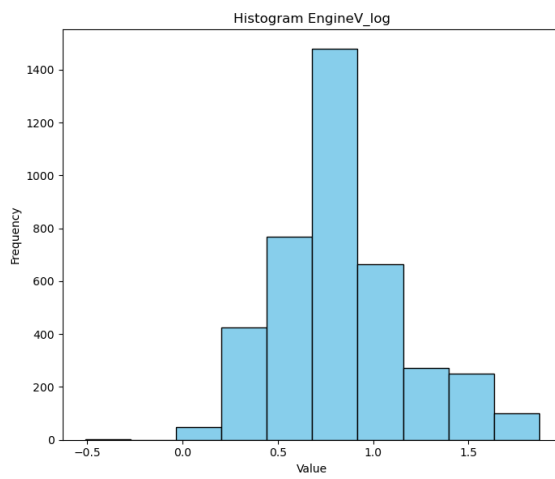

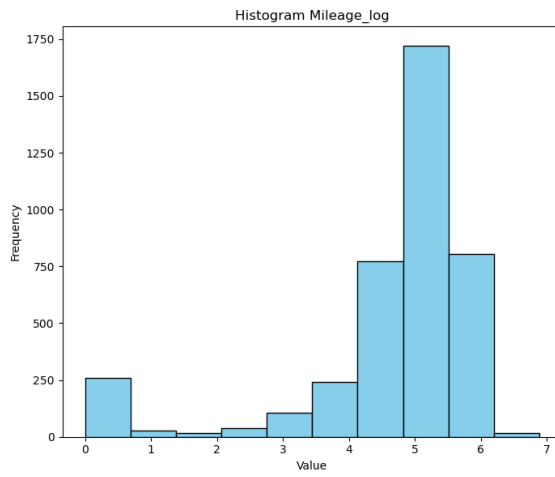


```
[12]:
```

	Price	Mileage	EngineV	Year
Price	1.000000	-0.473036	0.448590	0.485717
Mileage	-0.473036	1.000000	-0.034214	-0.664027
EngineV	0.448590	-0.034214	1.000000	0.038890
Year	0.485717	-0.664027	0.038890	1.000000

13 10. Data Transformation

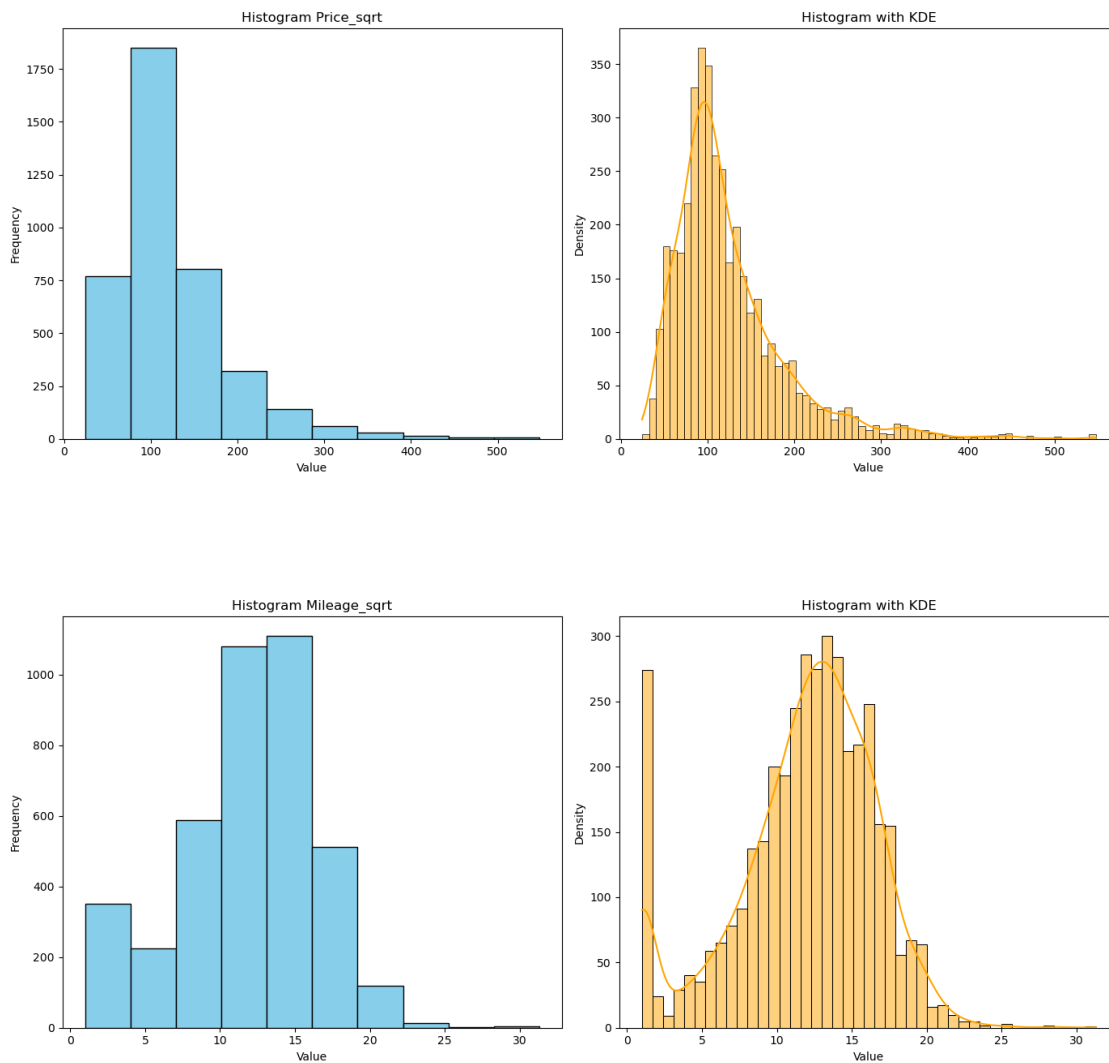
```
[13]: cars_price_df["Mileage_log"] = np.log(cars_price_df["Mileage"])
cars_price_df["Price_log"] = np.log(cars_price_df["Price"])
cars_price_df["EngineV_log"] = np.log(cars_price_df["EngineV"])
```


```
[14]: cars_price_num_df = cars_price_df[['Price_log', 'Mileage_log', 'EngineV_log']].
      ↪ copy()
      # Fit and transform your data using the pipeline
      processed_data = pipeline_num_var.fit_transform(cars_price_num_df)
```



```
[15]: cars_price_df["Mileage_sqrt"]=np.sqrt(cars_price_df["Mileage"])
cars_price_df["Price_sqrt"]=np.sqrt(cars_price_df["Price"])
cars_price_df["EngineV_sqrt"]=np.sqrt(cars_price_df["EngineV"])

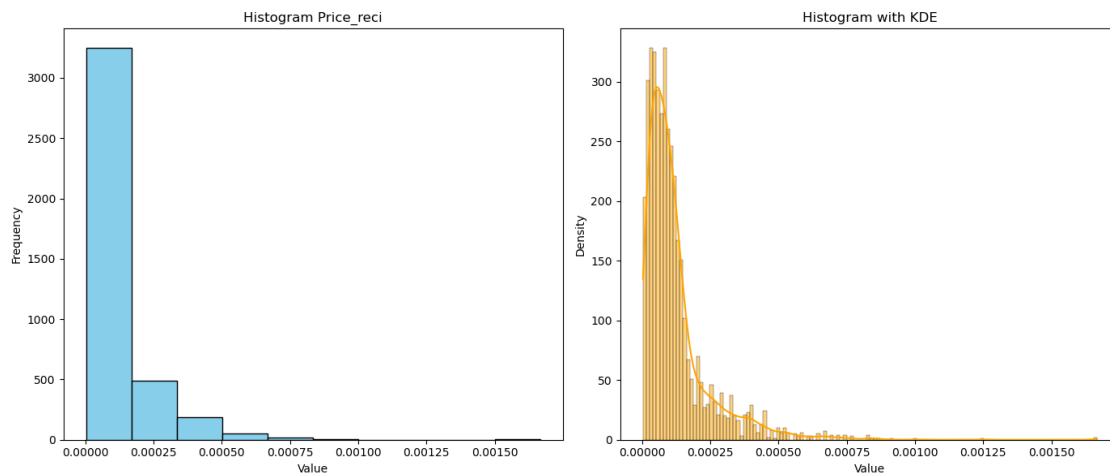
[16]: cars_price_num_df = cars_price_df[['Price_sqrt', 'Mileage_sqrt',
↳ 'EngineV_sqrt']].copy()
# Fit and transform your data using the pipeline
processed_data = pipeline_num_var.fit_transform(cars_price_num_df)
```

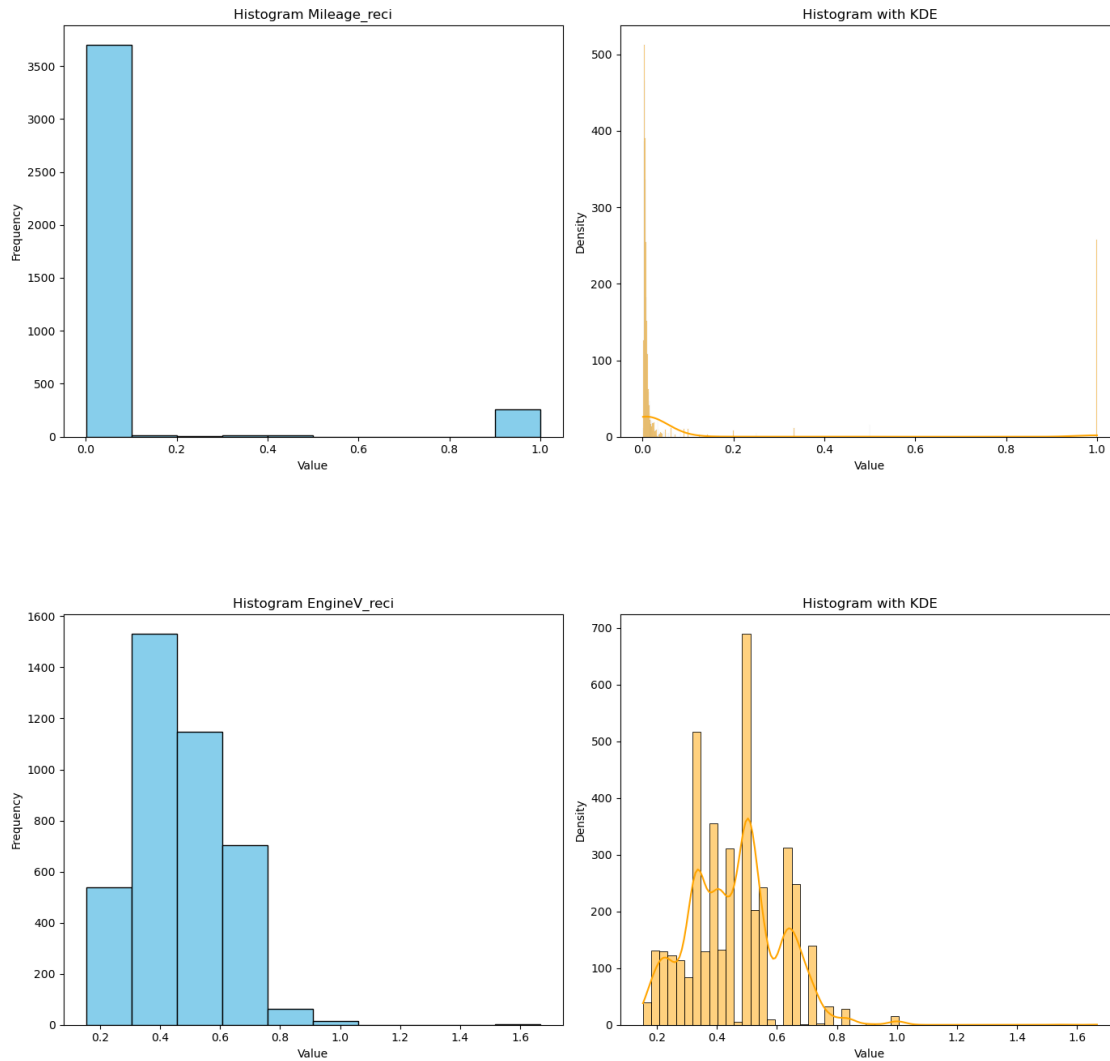




```
[17]: cars_price_df["Mileage_reci"]=1/(cars_price_df["Mileage"])
cars_price_df["Price_reci"]=1/(cars_price_df["Price"])
cars_price_df["EngineV_reci"]=1/(cars_price_df["EngineV"])
```

```
[18]: cars_price_num_df = cars_price_df[['Price_reci', 'Mileage_reci', 'EngineV_reci']].copy()
# Fit and transform your data using the pipeline
processed_data = pipeline_num_var.fit_transform(cars_price_num_df)
```





14 11. Standization - Normalization

```
[19]: cars_price_df[['Price_log', 'Mileage_sqrt', 'EngineV_reci']]
```

```
[19]:
```

	Price_log	Mileage_sqrt	EngineV_reci
0	12.310433	1.000000	0.158730
1	12.083905	1.000000	0.181818
2	12.088285	1.000000	0.181818
3	12.206068	1.000000	0.181818
4	12.206068	1.000000	0.181818
...
3998	6.396930	3.162278	0.666667
3999	8.003029	17.320508	0.357143
4000	7.740664	16.155494	0.434783

4001	8.612503	20.976177	0.500000
4002	10.463075	12.247449	0.357143

[4003 rows x 3 columns]

```
[20]: scaler = StandardScaler()

# Fit and transform the scaler on the selected columns
scaled_columns = scaler.
    ↪fit_transform(cars_price_df[['Price_log', 'Mileage_sqrt', 'EngineV_reciprocal']])

# Replace the original columns with the scaled columns
cars_price_df[['Price_log_scaler', 'Mileage_sqrt_scaler', 'EngineV_reciprocal_scaler']] = scaled_columns

print(cars_price_df)
```

	Brand	Price	Mileage	EngineV	Year	Mileage_log	Price_log	\
0	Mercedes-Benz	222000.0	1	6.3	2016	0.000000	12.310433	
1	Mercedes-Benz	177000.0	1	5.5	2016	0.000000	12.083905	
2	Mercedes-Benz	177777.0	1	5.5	2016	0.000000	12.088285	
3	Mercedes-Benz	199999.0	1	5.5	2016	0.000000	12.206068	
4	Mercedes-Benz	199999.0	1	5.5	2016	0.000000	12.206068	
...	
3998	Toyota	600.0	10	1.5	1979	2.302585	6.396930	
3999	Mercedes-Benz	2990.0	300	2.8	1979	5.703782	8.003029	
4000	Mercedes-Benz	2300.0	261	2.3	1978	5.564520	7.740664	
4001	Mercedes-Benz	5500.0	440	2.0	1978	6.086775	8.612503	
4002	Mercedes-Benz	34999.0	150	2.8	1969	5.010635	10.463075	

	EngineV_log	Mileage_sqrt	Price_sqrt	EngineV_sqrt	Mileage_reciprocal	\
0	1.840550	1.000000	471.168760	2.509980	1.000000	
1	1.704748	1.000000	420.713679	2.345208	1.000000	
2	1.704748	1.000000	421.636099	2.345208	1.000000	
3	1.704748	1.000000	447.212477	2.345208	1.000000	
4	1.704748	1.000000	447.212477	2.345208	1.000000	
...	
3998	0.405465	3.162278	24.494897	1.224745	0.100000	
3999	1.029619	17.320508	54.680892	1.673320	0.003333	
4000	0.832909	16.155494	47.958315	1.516575	0.003831	
4001	0.693147	20.976177	74.161985	1.414214	0.002273	
4002	1.029619	12.247449	187.080197	1.673320	0.006667	

	Price_reciprocal	EngineV_reciprocal	Price_log_scaler	Mileage_sqrt_scaler	\
0	0.000005	0.158730	3.104998	-2.301687	
1	0.000006	0.181818	2.862201	-2.301687	
2	0.000006	0.181818	2.866896	-2.301687	
3	0.000005	0.181818	2.993137	-2.301687	

4	0.000005	0.181818	2.993137	-2.301687
...
3998	0.001667	0.666667	-3.233208	-1.844164
3999	0.000334	0.357143	-1.511761	1.151619
4000	0.000435	0.434783	-1.792968	0.905111
4001	0.000182	0.500000	-0.858514	1.925134
4002	0.000029	0.357143	1.124964	0.078195

	EngineV_reciprocal_scaler
0	-2.000010
1	-1.845941
2	-1.845941
3	-1.845941
4	-1.845941
...	...
3998	1.389519
3999	-0.675975
4000	-0.157875
4001	0.277330
4002	-0.675975

[4003 rows x 17 columns]

```
[21]: # Initialize the MinMaxScaler
scaler = MinMaxScaler()

# Fit and transform the scaler on the selected column
scaled_column = scaler.fit_transform(cars_price_df[['Year']])

# Replace the original column with the scaled column
cars_price_df['Year_MinMax'] = scaled_column
print(cars_price_df)
```

	Brand	Price	Mileage	EngineV	Year	Mileage_log	Price_log	\
0	Mercedes-Benz	222000.0	1	6.3	2016	0.000000	12.310433	
1	Mercedes-Benz	177000.0	1	5.5	2016	0.000000	12.083905	
2	Mercedes-Benz	177777.0	1	5.5	2016	0.000000	12.088285	
3	Mercedes-Benz	199999.0	1	5.5	2016	0.000000	12.206068	
4	Mercedes-Benz	199999.0	1	5.5	2016	0.000000	12.206068	
...	
3998	Toyota	600.0	10	1.5	1979	2.302585	6.396930	
3999	Mercedes-Benz	2990.0	300	2.8	1979	5.703782	8.003029	
4000	Mercedes-Benz	2300.0	261	2.3	1978	5.564520	7.740664	
4001	Mercedes-Benz	5500.0	440	2.0	1978	6.086775	8.612503	
4002	Mercedes-Benz	34999.0	150	2.8	1969	5.010635	10.463075	

	EngineV_log	Mileage_sqrt	Price_sqrt	EngineV_sqrt	Mileage_reciprocal	\
0	1.840550	1.000000	471.168760	2.509980	1.000000	

1	1.704748	1.000000	420.713679	2.345208	1.000000
2	1.704748	1.000000	421.636099	2.345208	1.000000
3	1.704748	1.000000	447.212477	2.345208	1.000000
4	1.704748	1.000000	447.212477	2.345208	1.000000
...
3998	0.405465	3.162278	24.494897	1.224745	0.100000
3999	1.029619	17.320508	54.680892	1.673320	0.003333
4000	0.832909	16.155494	47.958315	1.516575	0.003831
4001	0.693147	20.976177	74.161985	1.414214	0.002273
4002	1.029619	12.247449	187.080197	1.673320	0.006667

	Price_reci	EngineV_reci	Price_log_scaler	Mileage_sqrt_scaler	\
0	0.000005	0.158730	3.104998	-2.301687	
1	0.000006	0.181818	2.862201	-2.301687	
2	0.000006	0.181818	2.866896	-2.301687	
3	0.000005	0.181818	2.993137	-2.301687	
4	0.000005	0.181818	2.993137	-2.301687	
...	
3998	0.001667	0.666667	-3.233208	-1.844164	
3999	0.000334	0.357143	-1.511761	1.151619	
4000	0.000435	0.434783	-1.792968	0.905111	
4001	0.000182	0.500000	-0.858514	1.925134	
4002	0.000029	0.357143	1.124964	0.078195	

	EngineV_reci_scaler	Year_MinMax
0	-2.000010	1.000000
1	-1.845941	1.000000
2	-1.845941	1.000000
3	-1.845941	1.000000
4	-1.845941	1.000000
...
3998	1.389519	0.212766
3999	-0.675975	0.212766
4000	-0.157875	0.191489
4001	0.277330	0.191489
4002	-0.675975	0.000000

[4003 rows x 18 columns]

```
[22]: cars_price_df[['Price_log_scaler',
↪ 'Mileage_sqrt_scaler', 'EngineV_reci_scaler', 'Year_MinMax']].describe()
```

	Price_log_scaler	Mileage_sqrt_scaler	EngineV_reci_scaler	Year_MinMax
count	4.003000e+03	4.003000e+03	4.003000e+03	4003.000000
mean	5.680082e-17	-5.680082e-17	-1.278018e-16	0.795653
std	1.000125e+00	1.000125e+00	1.000125e+00	0.142453
min	-3.233208e+00	-2.301687e+00	-2.032602e+00	0.000000

25%	-6.000318e-01	-5.059319e-01	-8.348595e-01	0.723404
50%	-6.794118e-02	1.464036e-01	-2.599454e-02	0.829787
75%	6.224600e-01	6.956878e-01	6.480596e-01	0.914894
max	3.427728e+00	4.110631e+00	8.062655e+00	1.000000

15 12. Convert Cat - to - Numerical Columns

```
[23]: cat_onehot_df = pd.get_dummies(cars_price_df['Brand'], prefix='Category',
    ↪drop_first=False)

# Concatenate the dummy variables with the original DataFrame
df = pd.concat([cars_price_df, cat_onehot_df], axis=1)
df
```

```
[23]:
```

	Brand	Price	Mileage	EngineV	Year	Mileage_log	Price_log	\
0	Mercedes-Benz	222000.0	1	6.3	2016	0.000000	12.310433	
1	Mercedes-Benz	177000.0	1	5.5	2016	0.000000	12.083905	
2	Mercedes-Benz	177777.0	1	5.5	2016	0.000000	12.088285	
3	Mercedes-Benz	199999.0	1	5.5	2016	0.000000	12.206068	
4	Mercedes-Benz	199999.0	1	5.5	2016	0.000000	12.206068	
...	
3998	Toyota	600.0	10	1.5	1979	2.302585	6.396930	
3999	Mercedes-Benz	2990.0	300	2.8	1979	5.703782	8.003029	
4000	Mercedes-Benz	2300.0	261	2.3	1978	5.564520	7.740664	
4001	Mercedes-Benz	5500.0	440	2.0	1978	6.086775	8.612503	
4002	Mercedes-Benz	34999.0	150	2.8	1969	5.010635	10.463075	

	EngineV_log	Mileage_sqrt	Price_sqrt	...	Mileage_sqrt_scaler	\
0	1.840550	1.000000	471.168760	...	-2.301687	
1	1.704748	1.000000	420.713679	...	-2.301687	
2	1.704748	1.000000	421.636099	...	-2.301687	
3	1.704748	1.000000	447.212477	...	-2.301687	
4	1.704748	1.000000	447.212477	...	-2.301687	
...	
3998	0.405465	3.162278	24.494897	...	-1.844164	
3999	1.029619	17.320508	54.680892	...	1.151619	
4000	0.832909	16.155494	47.958315	...	0.905111	
4001	0.693147	20.976177	74.161985	...	1.925134	
4002	1.029619	12.247449	187.080197	...	0.078195	

	EngineV_reci_scaler	Year_MinMax	Category_Audi	Category_BMW	\
0	-2.000010	1.000000	False	False	
1	-1.845941	1.000000	False	False	
2	-1.845941	1.000000	False	False	
3	-1.845941	1.000000	False	False	
4	-1.845941	1.000000	False	False	

...
3998	1.389519	0.212766	False	False
3999	-0.675975	0.212766	False	False
4000	-0.157875	0.191489	False	False
4001	0.277330	0.191489	False	False
4002	-0.675975	0.000000	False	False

	Category_Mercedes-Benz	Category_Mitsubishi	Category_Renault	\
0	True	False	False	
1	True	False	False	
2	True	False	False	
3	True	False	False	
4	True	False	False	

...
3998	False	False	False
3999	True	False	False
4000	True	False	False
4001	True	False	False
4002	True	False	False

	Category_Toyota	Category_Volkswagen
0	False	False
1	False	False
2	False	False
3	False	False
4	False	False

...
3998	True	False
3999	False	False
4000	False	False
4001	False	False
4002	False	False

[4003 rows x 25 columns]

15.1 VIF

```
[24]: from statsmodels.stats.outliers_influence import variance_inflation_factor
# VIF dataframe
# the independent variables set
X = cars_price_df[['Mileage_sqrt_scaler', 'EngineV_reciprocal_scaler', 'Year_MinMax']]

# VIF dataframe
vif_data = pd.DataFrame()
vif_data["feature"] = X.columns
# calculating VIF for each feature
vif_data["VIF"] = [variance_inflation_factor(X.values, i)
```

```

        for i in range(len(X.columns))]

print(vif_data)

```

```

           feature      VIF
0  Mileage_sqrt_scaler  1.013955
1  EngineV_reciprocal  1.001218
2           Year_MinMax  1.012732

```

```
[25]: df.columns
```

```
[25]: Index(['Brand', 'Price', 'Mileage', 'EngineV', 'Year', 'Mileage_log',
          'Price_log', 'EngineV_log', 'Mileage_sqrt', 'Price_sqrt',
          'EngineV_sqrt', 'Mileage_reciprocal', 'Price_reciprocal', 'EngineV_reciprocal',
          'Price_log_scaler', 'Mileage_sqrt_scaler', 'EngineV_reciprocal_scaler',
          'Year_MinMax', 'Category_Audi', 'Category_BMW',
          'Category_Mercedes-Benz', 'Category_Mitsubishi', 'Category_Renault',
          'Category_Toyota', 'Category_Volkswagen'],
          dtype='object')
```

```
[26]: df.info()
```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 4003 entries, 0 to 4002
Data columns (total 25 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   Brand                                4003 non-null   object
1   Price                                4003 non-null   float64
2   Mileage                              4003 non-null   int64
3   EngineV                              4003 non-null   float64
4   Year                                  4003 non-null   int64
5   Mileage_log                          4003 non-null   float64
6   Price_log                            4003 non-null   float64
7   EngineV_log                          4003 non-null   float64
8   Mileage_sqrt                         4003 non-null   float64
9   Price_sqrt                           4003 non-null   float64
10  EngineV_sqrt                         4003 non-null   float64
11  Mileage_reciprocal                   4003 non-null   float64
12  Price_reciprocal                     4003 non-null   float64
13  EngineV_reciprocal                   4003 non-null   float64
14  Price_log_scaler                     4003 non-null   float64
15  Mileage_sqrt_scaler                  4003 non-null   float64
16  EngineV_reciprocal_scaler             4003 non-null   float64
17  Year_MinMax                           4003 non-null   float64
18  Category_Audi                         4003 non-null   bool
19  Category_BMW                         4003 non-null   bool
20  Category_Mercedes-Benz                4003 non-null   bool

```



```

21 Category_Mitsubishi      4003 non-null    bool
22 Category_Renault        4003 non-null    bool
23 Category_Toyota         4003 non-null    bool
24 Category_Volkswagen     4003 non-null    bool
dtypes: bool(7), float64(15), int64(2), object(1)
memory usage: 590.4+ KB

```

16 13. Inferential statistics test

```

[27]: for i in ['Category_Audi', 'Category_BMW', 'Category_Mercedes-Benz',
             ↪ 'Category_Mitsubishi', 'Category_Renault',
             ↪ 'Category_Toyota', 'Category_Volkswagen']:
        df[i] = df[i].astype(int)

[28]: import statsmodels.api as sm
X = sm.add_constant(df[
    ↪ ['Mileage_sqrt_scaler', 'EngineV_reciprocal_scaler', 'Year_MinMax', 'Category_Audi',
    ↪ 'Category_BMW', 'Category_Mercedes-Benz', 'Category_Mitsubishi',
    ↪ 'Category_Renault', 'Category_Toyota', 'Category_Volkswagen']])

# Fit a linear regression model
model = sm.OLS(df['Price_log_scaler'], X)
results = model.fit()

# Print the regression results summary
print(results.summary())

# Calculate the R2 score
print(f"R2 score: {results.rsquared}")

```

```

                                OLS Regression Results
=====
Dep. Variable:      Price_log_scaler      R-squared:      0.802
Model:              OLS                   Adj. R-squared:  0.801
Method:             Least Squares         F-statistic:    1792.
Date:               Sat, 01 Jun 2024       Prob (F-statistic): 0.00
Time:               16:37:36               Log-Likelihood: -2443.4
No. Observations:   4003                  AIC:          4907.
Df Residuals:       3993                  BIC:          4970.
Df Model:           9
Covariance Type:    nonrobust
=====
=====
                                coef      std err          t      P>|t|      [0.025
0.975]
-----
-----

```

```

const                -3.1326      0.046    -68.103      0.000      -3.223
-3.042
Mileage_sqrt_scaler  -0.2231      0.009    -24.054      0.000      -0.241
-0.205
EngineV_rec1_scaler  -0.3421      0.008    -40.651      0.000      -0.359
-0.326
Year_MinMax          4.4511      0.065     68.088      0.000      4.323
4.579
Category_Audi         -0.3800      0.021    -18.270      0.000      -0.421
-0.339
Category_BMW          -0.2507      0.018    -14.088      0.000      -0.286
-0.216
Category_Mercedes-Benz -0.1987      0.017    -11.815      0.000      -0.232
-0.166
Category_Mitsubishi   -0.6544      0.024    -27.500      0.000      -0.701
-0.608
Category_Renault      -0.7840      0.024    -33.161      0.000      -0.830
-0.738
Category_Toyota       -0.3957      0.020    -19.948      0.000      -0.435
-0.357
Category_Volkswagen   -0.4691      0.017    -27.776      0.000      -0.502
-0.436
=====
Omnibus:              375.806    Durbin-Watson:      1.560
Prob(Omnibus):         0.000    Jarque-Bera (JB):   2464.651
Skew:                 -0.161    Prob(JB):           0.00
Kurtosis:             6.830    Cond. No.           3.84e+15
=====

```

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

[2] The smallest eigenvalue is 4.9e-28. This might indicate that there are strong multicollinearity problems or that the design matrix is singular.

R2 score: 0.8015206758411855

Observations - Durbin-Watson (1.560) - Since this is within the range of 1.5 and 2.5, we would consider autocorrelation not to be problematic in this regression model.

17 14. ML - Linear Regression Model

[29]:

```
df
```

[29]:

```

      Brand    Price  Mileage  EngineV  Year  Mileage_log  Price_log  \
0  Mercedes-Benz  222000.0        1      6.3   2016      0.000000  12.310433
1  Mercedes-Benz  177000.0        1      5.5   2016      0.000000  12.083905
2  Mercedes-Benz  177777.0        1      5.5   2016      0.000000  12.088285

```

3	Mercedes-Benz	199999.0	1	5.5	2016	0.000000	12.206068
4	Mercedes-Benz	199999.0	1	5.5	2016	0.000000	12.206068
...
3998	Toyota	600.0	10	1.5	1979	2.302585	6.396930
3999	Mercedes-Benz	2990.0	300	2.8	1979	5.703782	8.003029
4000	Mercedes-Benz	2300.0	261	2.3	1978	5.564520	7.740664
4001	Mercedes-Benz	5500.0	440	2.0	1978	6.086775	8.612503
4002	Mercedes-Benz	34999.0	150	2.8	1969	5.010635	10.463075

	EngineV_log	Mileage_sqrt	Price_sqrt	...	Mileage_sqrt_scaler	\
0	1.840550	1.000000	471.168760	...	-2.301687	
1	1.704748	1.000000	420.713679	...	-2.301687	
2	1.704748	1.000000	421.636099	...	-2.301687	
3	1.704748	1.000000	447.212477	...	-2.301687	
4	1.704748	1.000000	447.212477	...	-2.301687	
...	
3998	0.405465	3.162278	24.494897	...	-1.844164	
3999	1.029619	17.320508	54.680892	...	1.151619	
4000	0.832909	16.155494	47.958315	...	0.905111	
4001	0.693147	20.976177	74.161985	...	1.925134	
4002	1.029619	12.247449	187.080197	...	0.078195	

	EngineV_reciprocal_scaler	Year_MinMax	Category_Audi	Category_BMW	\
0	-2.000010	1.000000	0	0	
1	-1.845941	1.000000	0	0	
2	-1.845941	1.000000	0	0	
3	-1.845941	1.000000	0	0	
4	-1.845941	1.000000	0	0	
...	
3998	1.389519	0.212766	0	0	
3999	-0.675975	0.212766	0	0	
4000	-0.157875	0.191489	0	0	
4001	0.277330	0.191489	0	0	
4002	-0.675975	0.000000	0	0	

	Category_Mercedes-Benz	Category_Mitsubishi	Category_Renault	\
0	1	0	0	
1	1	0	0	
2	1	0	0	
3	1	0	0	
4	1	0	0	
...	
3998	0	0	0	
3999	1	0	0	
4000	1	0	0	
4001	1	0	0	
4002	1	0	0	

	Category_Toyota	Category_Volkswagen
0	0	0
1	0	0
2	0	0
3	0	0
4	0	0
...
3998	1	0
3999	0	0
4000	0	0
4001	0	0
4002	0	0

[4003 rows x 25 columns]

```
[30]: df.columns
```

```
[30]: Index(['Brand', 'Price', 'Mileage', 'EngineV', 'Year', 'Mileage_log',
          'Price_log', 'EngineV_log', 'Mileage_sqrt', 'Price_sqrt',
          'EngineV_sqrt', 'Mileage_reci', 'Price_reci', 'EngineV_reci',
          'Price_log_scaler', 'Mileage_sqrt_scaler', 'EngineV_reci_scaler',
          'Year_MinMax', 'Category_Audi', 'Category_BMW',
          'Category_Mercedes-Benz', 'Category_Mitsubishi', 'Category_Renault',
          'Category_Toyota', 'Category_Volkswagen'],
          dtype='object')
```

```
[31]: df_final=df[['Price_log_scaler', 'Mileage_sqrt_scaler',
          'EngineV_reci_scaler', 'Year_MinMax', 'Category_Audi',
          'Category_BMW', 'Category_Mercedes-Benz',
          'Category_Mitsubishi', 'Category_Renault', 'Category_Toyota',
          'Category_Volkswagen']].copy()
```

```
[32]: X=df_final.drop(['Price_log_scaler'],axis='columns')
      Y=df_final['Price_log_scaler']
```

```
[33]: X_train,X_test,y_train,y_test=train_test_split(X,Y,test_size=0.30)
      print("train data length:",len(X_train))
      print("test data length:",len(X_test))
      X.columns
```

```
train data length: 2802
test data length: 1201
```

```
[33]: Index(['Mileage_sqrt_scaler', 'EngineV_reci_scaler', 'Year_MinMax',
          'Category_Audi', 'Category_BMW', 'Category_Mercedes-Benz',
          'Category_Mitsubishi', 'Category_Renault', 'Category_Toyota',
          'Category_Volkswagen'],
```

```
dtype='object')
```

17.1 14.1 Linear Regression

```
[34]: def adjusted_r_squared(y_true, y_pred, n_samples, n_features):  
    """  
    Calculate the adjusted R-squared score.  
  
    Parameters:  
    - y_true: array-like, true target values  
    - y_pred: array-like, predicted target values  
    - n_samples: int, number of samples (observations)  
    - n_features: int, number of features (predictors)  
  
    Returns:  
    - adjusted R-squared score  
    """  
    from sklearn.metrics import r2_score  
  
    r_squared = r2_score(y_true, y_pred)  
    adjusted_r_squared = 1 - (1 - r_squared) * ((n_samples - 1) / (n_samples -  
↪n_features - 1))  
  
    return adjusted_r_squared
```

```
[35]: model_results = {}  
def regression_matrix(model ,X_train,X_test,y_test, model_name):  
    print("Model Name ",model_name)  
    y_pred = model.predict(X_test)  
    train_r2_score=round(model.score(X_train,y_train),3)  
    print("train R2 Score:",train_r2_score)  
    test_r2_score=round(model.score(X_test,y_test),3)  
    print("Test R2 Score:",test_r2_score)  
    print("Test R2 score:",r2_score(y_test,y_pred))  
  
    mse = round(mean_squared_error(y_test,y_pred),3)  
    print("MSE:",mse)  
    #rmse=round(root_mean_squared_error(y_test,y_pred),3)  
    rmse=round(np.sqrt(mse),3)  
    print("RMSE:",rmse)  
    adj_r2_score=round(adjusted_r_squared(y_test,y_pred,len(y_test),len(X_train.  
↪columns)),3)  
    print("Adj-R Score",adj_r2_score)  
  
    if abs(train_r2_score - test_r2_score) > .10:  
        print("model : " , model_name , "is overfitting")  
    if train_r2_score < 0.50:
```

```

        print("model :" , model_name , "is underfitting")

    ↵
    ↪model_results[model_name]=[train_r2_score,test_r2_score,adj_r2_score,mse,rmse]

```

```

[36]: lr = LinearRegression()
      lr.fit(X_train,y_train)

      print("columns:",X_train.columns)
      print('Coefficients: ', lr.coef_)
      print('Intercept:',lr.intercept_)

      regression_matrix(lr ,X_train,X_test,y_test, "LinearReg")

columns: Index(['Mileage_sqrt_scaler', 'EngineV_reciprocal_scaler', 'Year_MinMax',
               'Category_Audi', 'Category_BMW', 'Category_Mercedes-Benz',
               'Category_Mitsubishi', 'Category_Renault', 'Category_Toyota',
               'Category_Volkswagen'],
              dtype='object')
Coefficients: [-2.26901611e-01 -3.34514321e-01  4.54020890e+00 -1.46260246e+13
               -1.46260246e+13 -1.46260246e+13 -1.46260246e+13 -1.46260246e+13
               -1.46260246e+13 -1.46260246e+13]
Intercept: 14626024585324.809
Model Name  LinearReg
train R2 Score: 0.807
Test R2 Score: 0.785
Test R2 score: 0.7846599620798789
MSE: 0.197
RMSE: 0.444
Adj-R Score 0.783

```

17.1.1 Lasso Regression - L1

```

[37]: lasso_reg = Lasso(alpha=0.1) # Regularization strength (alpha) is set to 0.1
      lasso_reg.fit(X_train,y_train)

      print("columns:",X_train.columns)
      print('Coefficients: ', lasso_reg.coef_)
      print('Intercept:',lasso_reg.intercept_)

      regression_matrix(lasso_reg ,X_train,X_test,y_test, "Lasso")

columns: Index(['Mileage_sqrt_scaler', 'EngineV_reciprocal_scaler', 'Year_MinMax',
               'Category_Audi', 'Category_BMW', 'Category_Mercedes-Benz',
               'Category_Mitsubishi', 'Category_Renault', 'Category_Toyota',

```

```

        'Category_Volkswagen'],
        dtype='object')
Coefficients: [-0.52324899 -0.33562024  0.          -0.          0.          0.
 -0.          -0.          0.          -0.          ]
Intercept: -0.014651228818764964
Model Name Lasso
train R2 Score: 0.535
Test R2 Score: 0.546
Test R2 score: 0.5464730454365776
MSE: 0.416
RMSE: 0.645
Adj-R Score 0.543

```

17.1.2 Ridge Regression -L2

```
[38]: ridge_reg = Ridge(alpha=0.1) # Regularization strength (alpha) is set to 0.1
ridge_reg.fit(X_train,y_train)
```

```

print("columns:",X_train.columns)
print('Coefficients: ', ridge_reg.coef_)
print('Intercept:',ridge_reg.intercept_)

regression_matrix(ridge_reg ,X_train,X_test,y_test, "Ridge")

```

```

columns: Index(['Mileage_sqrt_scaler', 'EngineV_reciprocal_scaler', 'Year_MinMax',
               'Category_Audi', 'Category_BMW', 'Category_Mercedes-Benz',
               'Category_Mitsubishi', 'Category_Renault', 'Category_Toyota',
               'Category_Volkswagen'],
              dtype='object')
Coefficients: [-0.22962844 -0.33578971  4.50788691  0.06821927  0.2005005
 0.2554526
 -0.20342715 -0.34470256  0.04436584 -0.0204085 ]
Intercept: -3.6321460198717666
Model Name Ridge
train R2 Score: 0.807
Test R2 Score: 0.785
Test R2 score: 0.784952991018684
MSE: 0.197
RMSE: 0.444
Adj-R Score 0.783

```

17.1.3 Elastic Net

```
[39]: elastic_net = ElasticNet(alpha=0.1, l1_ratio=0.5) # l1_ratio controls the
↪balance between L1 and L2 penalties
elastic_net.fit(X_train,y_train)
```

```
print("columns:",X_train.columns)
print('Coefficients: ', elastic_net.coef_)
print('Intercept:',elastic_net.intercept_)
```

```
regression_matrix(elastic_net ,X_train,X_test,y_test, "ElasticNet")
```

```
columns: Index(['Mileage_sqrt_scaler', 'EngineV_reciprocal_scaler', 'Year_MinMax',
               'Category_Audi', 'Category_BMW', 'Category_Mercedes-Benz',
               'Category_Mitsubishi', 'Category_Renault', 'Category_Toyota',
               'Category_Volkswagen'],
              dtype='object')
Coefficients: [-0.53140875 -0.36702749  0.17807115 -0.          0.          0.
               -0.          -0.          0.          -0.          ]
Intercept: -0.15596468992061246
Model Name ElasticNet
train R2 Score: 0.562
Test R2 Score: 0.572
Test R2 score: 0.5723694410669489
MSE: 0.392
RMSE: 0.626
Adj-R Score 0.569
```

17.2 14.2 Decision Tree Regression

```
[40]: # Create and fit the model
model = DecisionTreeRegressor()
model.fit(X_train,y_train)
print("Model - Decision Tree Regression")

regression_matrix(model ,X_train,X_test,y_test, "DT")
```

```
Model - Decision Tree Regression
Model Name DT
train R2 Score: 0.996
Test R2 Score: 0.765
Test R2 score: 0.765061007871335
MSE: 0.215
RMSE: 0.464
Adj-R Score 0.763
model : DT is overfitting
```

17.3 14.3 Random Forest Regression

```
[41]: # Create and fit the model
model = RandomForestRegressor()
model.fit(X_train, y_train)
print("Model - Random Forest Regression")
```



```
regression_matrix(model ,X_train,X_test,y_test, "RandomForest")
```

Model - Random Forest Regression
Model Name RandomForest
train R2 Score: 0.978
Test R2 Score: 0.855
Test R2 score: 0.8550559993534799
MSE: 0.133
RMSE: 0.365
Adj-R Score 0.854
model : RandomForest is overfitting

17.4 14.4 Support Vector Regression (SVR)

```
[42]: # Create and fit the model
model = SVR(kernel='linear')
model.fit(X_train, y_train)
print("Model - Support Vector Regression ")

regression_matrix(model ,X_train,X_test,y_test, "SVR")
```

Model - Support Vector Regression
Model Name SVR
train R2 Score: 0.803
Test R2 Score: 0.784
Test R2 score: 0.7840000160244213
MSE: 0.198
RMSE: 0.445
Adj-R Score 0.782

17.5 14.5 AdaBoost Regression

```
[43]: # Create and fit the model
ada_boost = AdaBoostRegressor()
ada_boost.fit(X_train, y_train)
print("Model - AdaBoost Regression ")

regression_matrix(ada_boost ,X_train,X_test,y_test, "AdaBoost")
```

Model - AdaBoost Regression
Model Name AdaBoost
train R2 Score: 0.812
Test R2 Score: 0.774
Test R2 score: 0.7742414400136288
MSE: 0.207
RMSE: 0.455
Adj-R Score 0.772

17.6 14.6 Gradient Boosting Regression

```
[44]: # Create and fit the model
gradient_boost = GradientBoostingRegressor()
gradient_boost.fit(X_train, y_train)
print("Model - Gradient Boosting Regression")

regression_matrix(gradient_boost ,X_train,X_test,y_test, "GradientBoost")
```

Model - Gradient Boosting Regression
Model Name GradientBoost
train R2 Score: 0.891
Test R2 Score: 0.858
Test R2 score: 0.857631467748983
MSE: 0.13
RMSE: 0.361
Adj-R Score 0.856

17.7 14.7 XGBoost Regression

```
[45]: # Create and fit the model
xg_boost = xgb.XGBRegressor()
xg_boost.fit(X_train, y_train)
print("Model-XGBoost Regression")

regression_matrix(xg_boost ,X_train,X_test,y_test, "XGB")
```

Model-XGBoost Regression
Model Name XGB
train R2 Score: 0.971
Test R2 Score: 0.85
Test R2 score: 0.8499927128937329
MSE: 0.137
RMSE: 0.37
Adj-R Score 0.849
model : XGB is overfitting

17.8 18. Summary

```
[46]: print("\n\n")
result=pd.DataFrame(model_results,index=["Train R2","Test R2" ,"Adj_R2","MSE","RMSE"])
print(result)
print("\n\n")

print(tabulate(result, headers='keys', tablefmt='pretty'))
```

	LinearReg	Lasso	Ridge	ElasticNet	DT	RandomForest	SVR	\
Train R2	0.807	0.535	0.807	0.562	0.996	0.978	0.803	
Test R2	0.785	0.546	0.785	0.572	0.765	0.855	0.784	
Adj R2	0.783	0.543	0.783	0.569	0.763	0.854	0.782	
MSE	0.197	0.416	0.197	0.392	0.215	0.133	0.198	
RMSE	0.444	0.645	0.444	0.626	0.464	0.365	0.445	

	AdaBoost	GradientBoost	XGB
Train R2	0.812	0.891	0.971
Test R2	0.774	0.858	0.850
Adj R2	0.772	0.856	0.849
MSE	0.207	0.130	0.137
RMSE	0.455	0.361	0.370

```

+-----+-----+-----+-----+-----+-----+-----+-----+
--+-----+-----+-----+-----+
|           | LinearReg | Lasso | Ridge | ElasticNet | DT   | RandomForest | SVR |
+-----+-----+-----+-----+-----+-----+-----+-----+
--+-----+-----+-----+-----+
| Train R2 | 0.807 | 0.535 | 0.807 | 0.562 | 0.996 | 0.978 | 0.803 |
| Test R2  | 0.785 | 0.546 | 0.785 | 0.572 | 0.765 | 0.855 | 0.784 |
| Adj R2   | 0.783 | 0.543 | 0.783 | 0.569 | 0.763 | 0.854 | 0.782 |
| MSE      | 0.197 | 0.416 | 0.197 | 0.392 | 0.215 | 0.133 | 0.198 |
| RMSE     | 0.444 | 0.645 | 0.444 | 0.626 | 0.464 | 0.365 | 0.445 |
+-----+-----+-----+-----+-----+-----+-----+-----+
--+-----+-----+-----+-----+

```

[]: