


Iris_Flowers_Classification

May 30, 2024

- 1 Objective : Iris Flowers Classification
- 2 Exploratory Data Analysis (EDA) - Python
- 3 Insights - Patterns
- 4 Classification (Using the ML)



Iris Versicolor



Iris Setosa



Iris Virginica

5 1. Load Python Modules

```
[1]: # Use Python's import statement to load modules
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns

from tabulate import tabulate

from sklearn.base import BaseEstimator, TransformerMixin
from sklearn.pipeline import Pipeline
from sklearn.preprocessing import LabelEncoder
```

```

from sklearn.preprocessing import StandardScaler
from sklearn.preprocessing import MinMaxScaler
from sklearn.model_selection import train_test_split

from sklearn.linear_model import LogisticRegression

from sklearn.naive_bayes import CategoricalNB
from sklearn.naive_bayes import GaussianNB

from sklearn.svm import SVC

from sklearn.tree import DecisionTreeClassifier
from sklearn.ensemble import RandomForestClassifier

from sklearn.ensemble import AdaBoostClassifier
from sklearn.tree import DecisionTreeClassifier

from sklearn.ensemble import GradientBoostingClassifier
from xgboost import XGBClassifier

from sklearn.metrics import accuracy_score
from sklearn.metrics import classification_report
from sklearn.metrics import confusion_matrix

```

6 2. Read the Dataset from CSV file - Using Pandas

```

[2]: file_path=r"iris_dataset.csv"
iris_df=pd.read_csv(file_path)
iris_df

```

```

[2]:      sepal length (cm)  sepal width (cm)  petal length (cm)  petal width (cm)  \
0                5.1             3.5             1.4             0.2
1                4.9             3.0             1.4             0.2
2                4.7             3.2             1.3             0.2
3                4.6             3.1             1.5             0.2
4                5.0             3.6             1.4             0.2
..                ...             ...             ...             ...
145              6.7             3.0             5.2             2.3
146              6.3             2.5             5.0             1.9
147              6.5             3.0             5.2             2.0
148              6.2             3.4             5.4             2.3
149              5.9             3.0             5.1             1.8

      target target_names
0          0      setosa

```

```

1         0      setosa
2         0      setosa
3         0      setosa
4         0      setosa
..      ...      ...
145      2     virginica
146      2     virginica
147      2     virginica
148      2     virginica
149      2     virginica

```

[150 rows x 6 columns]

```

[3]: #rename the columns names
iris_df.columns = ['sepal_length', 'sepal_width', 'petal_length', 'petal_width', 'species_labels', 'species']

```

7 3. Basic Inspection on given dataset

```

[4]: def basic_inspection_dataset(table):
    print("Top 5 Records of dataset")
    print(table.head())
    print()

    print("Bottom Records of dataset")
    print(table.tail())
    print()

    print("Column/features/Variable - Names of Given dataset")
    print(table.columns)
    print()

    print("Shape(rows x columns) - of Given dataset")
    print(table.shape)
    print()

    print("Data types - Given Column Names")
    print(table.dtypes)
    print()

    print("Summry of dataset")
    print(table.info())
    print()

    print("To see the count of null/nan values in columns of dataset")
    print(table.isnull().value_counts())

```

```

print()

print("Dataset Summary ")
print(table.describe())
print()

basic_inspection_dataset(iris_df)

```

Top 5 Records of dataset

	sepal_length	sepal_width	petal_length	petal_width	species_labels	\
0	5.1	3.5	1.4	0.2	0	
1	4.9	3.0	1.4	0.2	0	
2	4.7	3.2	1.3	0.2	0	
3	4.6	3.1	1.5	0.2	0	
4	5.0	3.6	1.4	0.2	0	

```

species
0 setosa
1 setosa
2 setosa
3 setosa
4 setosa

```

Bottom Records of dataset

	sepal_length	sepal_width	petal_length	petal_width	species_labels	\
145	6.7	3.0	5.2	2.3	2	
146	6.3	2.5	5.0	1.9	2	
147	6.5	3.0	5.2	2.0	2	
148	6.2	3.4	5.4	2.3	2	
149	5.9	3.0	5.1	1.8	2	

```

species
145 virginica
146 virginica
147 virginica
148 virginica
149 virginica

```

Column/features/Variable - Names of Given dataset

```

Index(['sepal_length', 'sepal_width', 'petal_length', 'petal_width',
      'species_labels', 'species'],
      dtype='object')

```

Shape(rows x columns) - of Given dataset
(150, 6)

Data types - Given Column Names

```

sepal_length      float64

```

```

sepal_width      float64
petal_length     float64
petal_width      float64
species_labels   int64
species          object
dtype: object

```

Summary of dataset

```
<class 'pandas.core.frame.DataFrame'>
```

RangeIndex: 150 entries, 0 to 149

Data columns (total 6 columns):

#	Column	Non-Null Count	Dtype
0	sepal_length	150 non-null	float64
1	sepal_width	150 non-null	float64
2	petal_length	150 non-null	float64
3	petal_width	150 non-null	float64
4	species_labels	150 non-null	int64
5	species	150 non-null	object

dtypes: float64(4), int64(1), object(1)

memory usage: 7.2+ KB

None

To see the count of null/nan values in columns of dataset

sepal_length	sepal_width	petal_length	petal_width	species_labels	species
False	False	False	False	False	False
150					

Name: count, dtype: int64

Dataset Summary

	sepal_length	sepal_width	petal_length	petal_width	species_labels
count	150.000000	150.000000	150.000000	150.000000	150.000000
mean	5.843333	3.057333	3.758000	1.199333	1.000000
std	0.828066	0.435866	1.765298	0.762238	0.819232
min	4.300000	2.000000	1.000000	0.100000	0.000000
25%	5.100000	2.800000	1.600000	0.300000	0.000000
50%	5.800000	3.000000	4.350000	1.300000	1.000000
75%	6.400000	3.300000	5.100000	1.800000	2.000000
max	7.900000	4.400000	6.900000	2.500000	2.000000

7.1 Observations on Iris Dataset

1. Dataset Overview

- Number of Records: 150
- Features/Variables: 5
- Independent Features/Variables: sepal_length, sepal_width, petal_length, petal_width (Numerical/Continuous)

- Dependent Feature/Variable: species (Categorical)
- 2. Data Integrity**
 - No missing values are observed in any of the columns.
 - 3. Data Types**
 - Independent features' data types: float/Real numbers
 - Output variable: Categorical Variable
 - 4. Summary Statistics**
 - **sepal_length**: min: 4.3, 25%: 5.1, mean: 5.8, 50%: 5.8, 75%: 6.4, max: 7.9,
 - **sepal_width**: min: 2.0, 25%: 2.8, mean: 3.1, 50%: 3.0, 75%: 3.3, max: 4.4,
 - **petal_length**: min: 1.0, 25%: 1.6, mean: 3.8, 50%: 4.3, 75%: 5.1, max: 6.9,
 - **petal_width**: min: 0.1, 25%: 0.3, mean: 1.2, 50%: 1.3, 75%: 1.8, max: 2.5,
 - 5. Observations on Spread and Range**
 - The spread is more pronounced in petal_length.
 - The range is broader in petal_length compared to other features.
 - 6. Lowest Mean/Median**
 - The lowest mean and median are observed in petal_width.
 - 7. Petal Width Distribution**
 - The range (25% - 75%) for petal_width is between 0.3 to 1.8 cm, with a median of 1.3 cm.
 - 8. Petal Length Distribution**
 - The range (25% - 75%) for petal_length is between 1.6 to 5.1 cm, with a median of 4.3 cm.
 - 9. Sepal Width Distribution**
 - The range (25% - 75%) for sepal_width is between 2.8 to 3.3 cm, with a median of 3.0 cm.
 - 10. Sepal Length Distribution**
 - The range (25% - 75%) for sepal_length is between 5.1 to 6.4 cm, with a median of 5.8 cm.

8 4. Handling Missing Values - Cat - Variables

```
[5]: iris_df.isnull().sum()
```

```
[5]: sepal_length      0
     sepal_width       0
     petal_length      0
     petal_width       0
     species_labels    0
     species          0
     dtype: int64
```

9 5. Categorical- UniVariable - Analysis -Using Pipeline

```
[6]: class BarPieChartTransformer(BaseEstimator, TransformerMixin):
    def __init__(self):
        pass

    def fit(self, X, y=None):
        return self

    def transform(self, X):
        df=X.copy()
        # get cat columns
        cat_cols = df.select_dtypes(include='object').columns
        for cat_name in cat_cols:
            value_counts = df[cat_name].value_counts().reset_index()
            # Rename the columns
            value_counts.columns = ['Class', 'Frequency']

            # Print the result as a table
            print(f"{cat_name} frequency table")
            print(tabulate(value_counts, headers='keys', tablefmt='pretty'))

            # Calculate relative frequency
            total_count = value_counts['Frequency'].sum()
            value_counts['Relative Frequency %'] =
↪round((value_counts['Frequency'] / total_count)*100,2)

            # Print the result as a table
            print(f"{cat_name} Relative frequency table")
            print(tabulate(value_counts, headers='keys', tablefmt='pretty'))

            # Extract the values and index from value counts
            value_counts = df[cat_name].value_counts()
            values = value_counts.values
            labels = value_counts.index

            fig, axs = plt.subplots(1, 2, figsize=(12, 6)) # 1 row, 2 columns
            # Create a bar graph
            axs[0].bar(labels, values)
            axs[0].set_title(f'Frequency of {cat_name}')
            axs[0].set_xlabel('Categories') # Set x-label
            axs[0].set_ylabel('Count') # Set y-label

            axs[1].pie(value_counts.values, labels=value_counts.index,
↪autopct='%1.1f%%', startangle=140)
            axs[1].set_title(f'Relative Frequency of {cat_name}')
            plt.tight_layout()
```

```
# Show the plot
plt.show()
```

```
[7]: pipeline_cat_var = Pipeline([
      ('cat_univariate_analysis', BarPieChartTransformer())
    ])


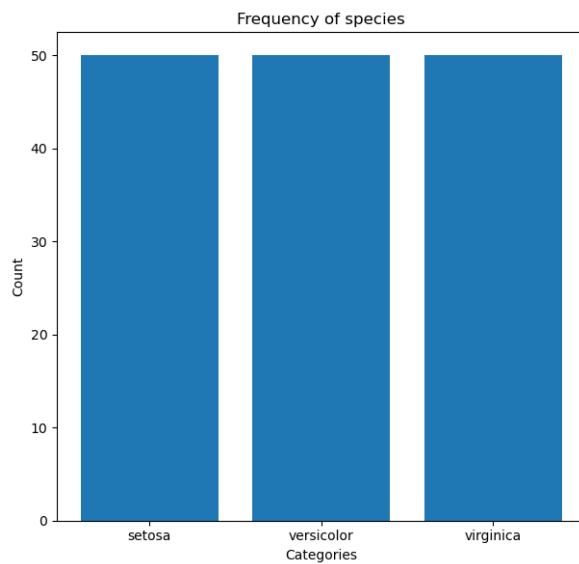
# Fit and transform your data using the pipeline
processed_data = pipeline_cat_var.fit_transform(iris_df)
```

species frequency table

	Class	Frequency
0	setosa	50
1	versicolor	50
2	virginica	50

species Relative frequency table

	Class	Frequency	Relative Frequency %
0	setosa	50	33.33
1	versicolor	50	33.33
2	virginica	50	33.33



10 6. Handling Missing Values in Numerical Columns

```
[8]: iris_df.isnull().sum()
```

```
[8]: sepal_length    0
     sepal_width    0
     petal_length   0
     petal_width    0
     species_labels  0
     species        0
     dtype: int64
```

```
[9]: iris_df.describe()
```

```
[9]:
```

	sepal_length	sepal_width	petal_length	petal_width	species_labels
count	150.000000	150.000000	150.000000	150.000000	150.000000
mean	5.843333	3.057333	3.758000	1.199333	1.000000
std	0.828066	0.435866	1.765298	0.762238	0.819232
min	4.300000	2.000000	1.000000	0.100000	0.000000
25%	5.100000	2.800000	1.600000	0.300000	0.000000
50%	5.800000	3.000000	4.350000	1.300000	1.000000
75%	6.400000	3.300000	5.100000	1.800000	2.000000
max	7.900000	4.400000	6.900000	2.500000	2.000000

11 7. Numerical - UniVariable - Analysis - Using -Pipeline

```
[10]: class HistBoxChartTransformer(BaseEstimator, TransformerMixin):
        def __init__(self):
            pass

        def fit(self, X, y=None):
            return self

        def transform(self, X):
            df=X.copy()
            # getting num cols
            num_cols = df.select_dtypes(exclude='object').columns
            for con_var in num_cols:

                # Create a figure and axes object
                fig, axes = plt.subplots(1, 2, figsize=(14, 6))

                # Plot histogram without KDE on the left
                axes[0].hist(df[con_var], color='skyblue', edgecolor='black')
                axes[0].set_xlabel('Value')
                axes[0].set_ylabel('Frequency')
```

```

axes[0].set_title(f'Histogram {con_var}')

# Plot histogram with KDE on the right
sns.histplot(data=df, x=con_var, kde=True, color='orange',
edgecolor='black', ax=axes[1])
axes[1].set_xlabel('Value')
axes[1].set_ylabel('Density')
axes[1].set_title('Histogram with KDE')

# Adjust layout
plt.tight_layout()

# Show the combined plot
plt.show()


```

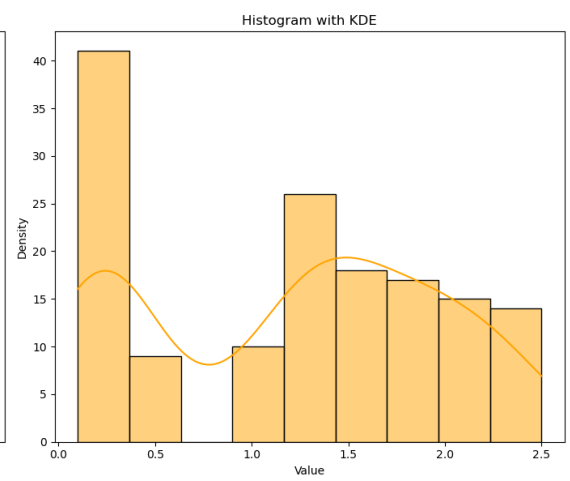
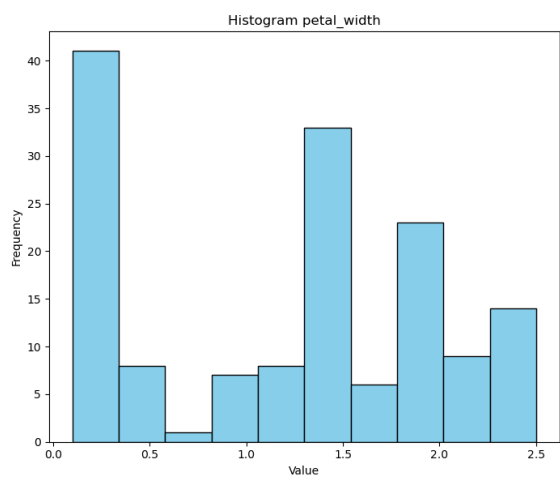


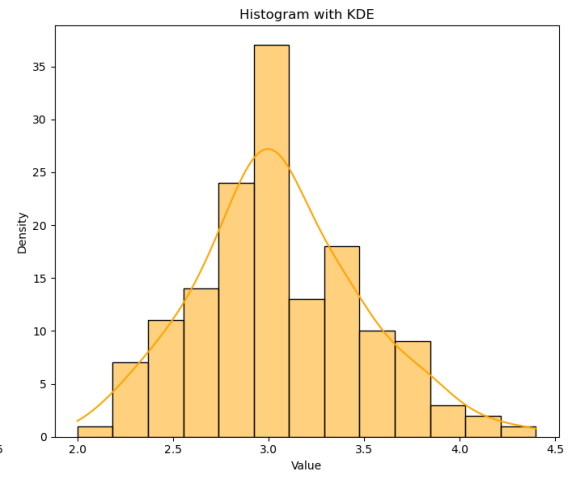
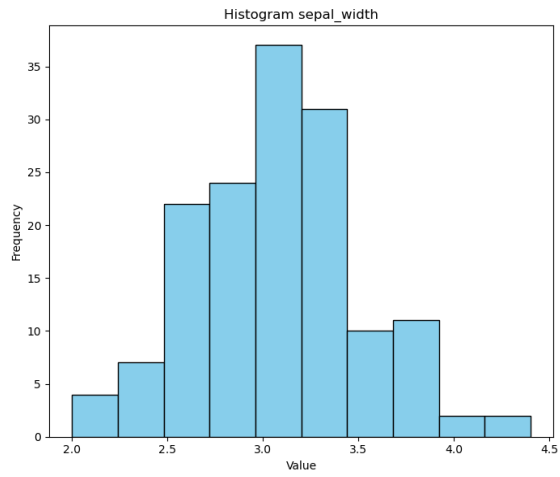
```

[11]: pipeline_num_var = Pipeline([
      ('num_uni_variate_analysis', HistBoxChartTransformer())
    ])

iris_num_df = iris_df[['sepal_length', 'sepal_width', 'petal_length',
      ↪ 'petal_width']].copy()
# Fit and transform your data using the pipeline
processed_data = pipeline_num_var.fit_transform(iris_num_df)

```





12 8. Numerical - Variables -Outliers Analysis

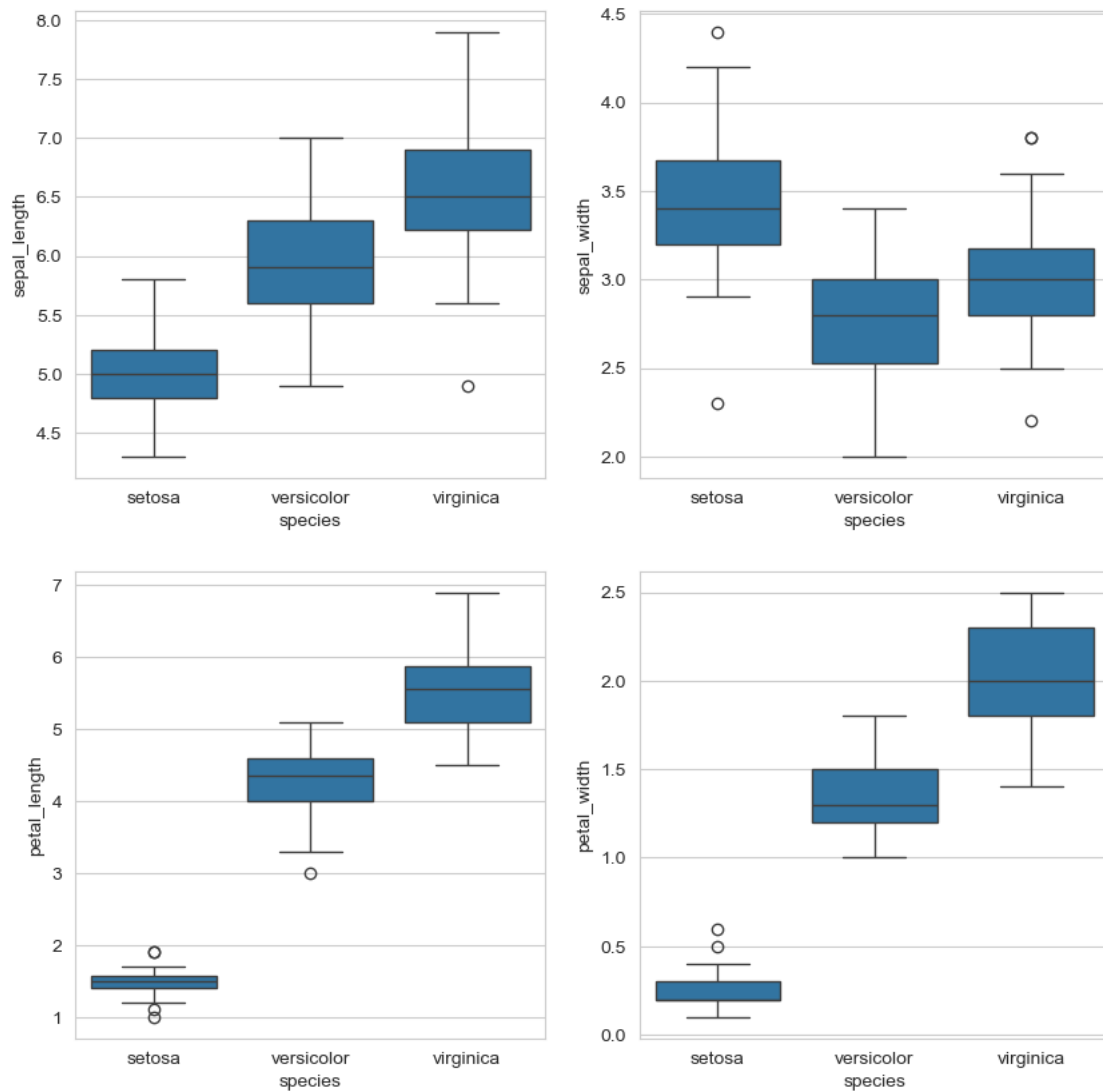
13 9. Bi Variate Analysis

13.1 9.1 Num Vs cat(target)

```
[12]: output_var='species'
```

```
[13]: sns.set_style("whitegrid")
fig, axes = plt.subplots(2, 2, figsize=(10, 10))
fig.suptitle('Box-Plots Features Vs Flower Type')
sns.boxplot(ax=axes[0, 0], x=output_var, y='sepal_length', data=iris_df)
sns.boxplot(ax=axes[0, 1], x=output_var, y='sepal_width', data=iris_df)
sns.boxplot(ax=axes[1, 0], x=output_var, y='petal_length', data=iris_df)
sns.boxplot(ax=axes[1, 1], x=output_var, y='petal_width', data=iris_df)
plt.show()
```

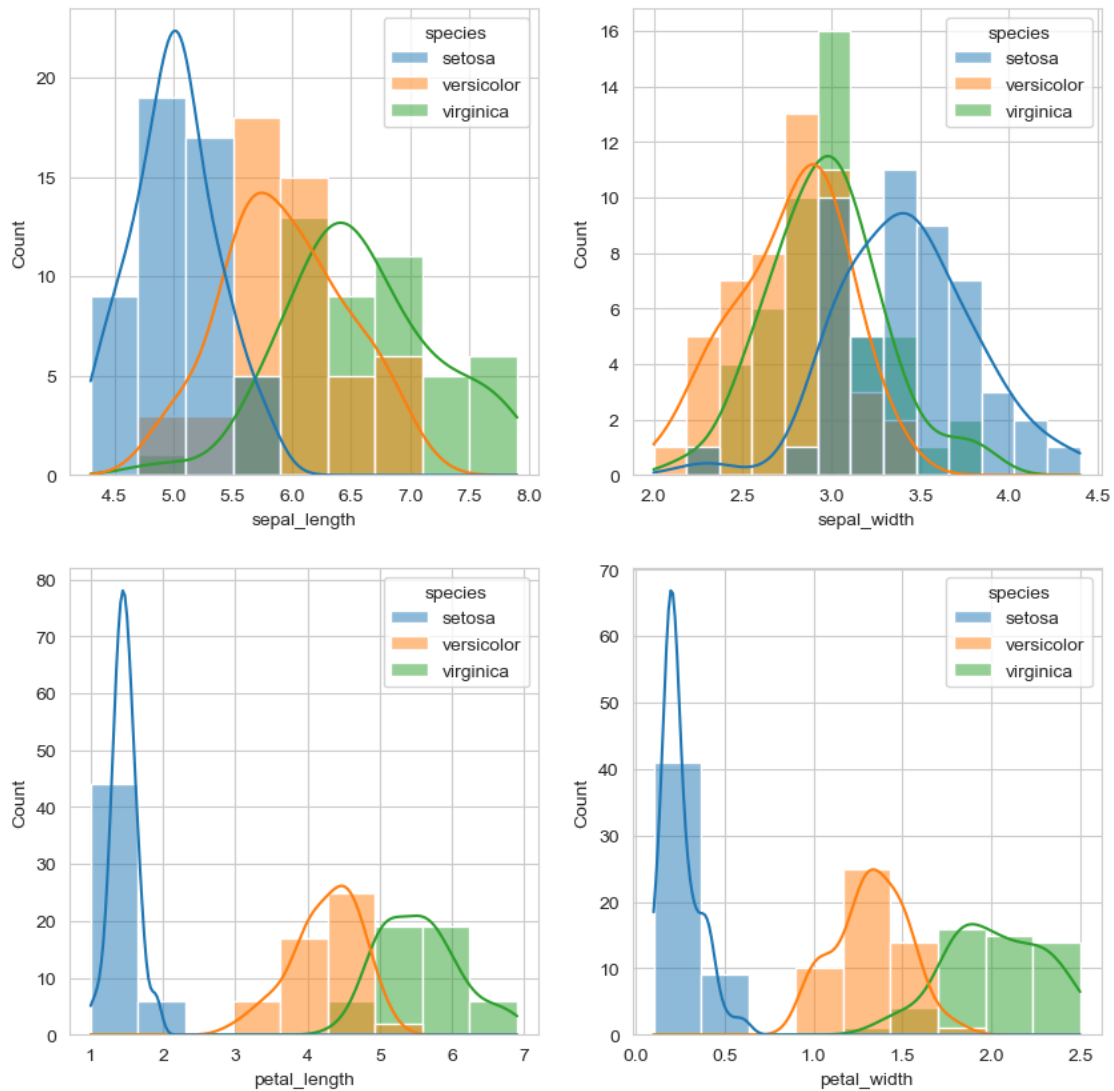
Box-Plots Features Vs Flower Type



```
[14]: sns.set_style("whitegrid")
fig, axes = plt.subplots(2, 2, figsize=(10, 10))
fig.suptitle('Kde-Plots')
sns.histplot(ax=axes[0, 0], hue=output_var, x='sepal_length',
             data=iris_df, kde=True)
sns.histplot(ax=axes[0, 1], hue=output_var, x='sepal_width',
             data=iris_df, kde=True)
sns.histplot(ax=axes[1, 0], hue=output_var, x='petal_length',
             data=iris_df, kde=True)
```

```
sns.histplot(ax=axes[1, 1], hue=output_var, x='petal_width',  
             data=iris_df, kde=True)  
plt.show()
```

Kde-Plots



13.2 9.2 Num Vs Num

```
[15]: # Selecting only numerical columns  
numerical_columns = ['sepal_length', 'sepal_width', 'petal_length',  
                    'petal_width']
```

```

# Creating unique scatter plots
num_cols_count = len(numerical_columns)
num_plots = num_cols_count * (num_cols_count - 1) // 2

# Setting up subplots
fig, axes = plt.subplots(num_plots // 2, 2, figsize=(15, 15))

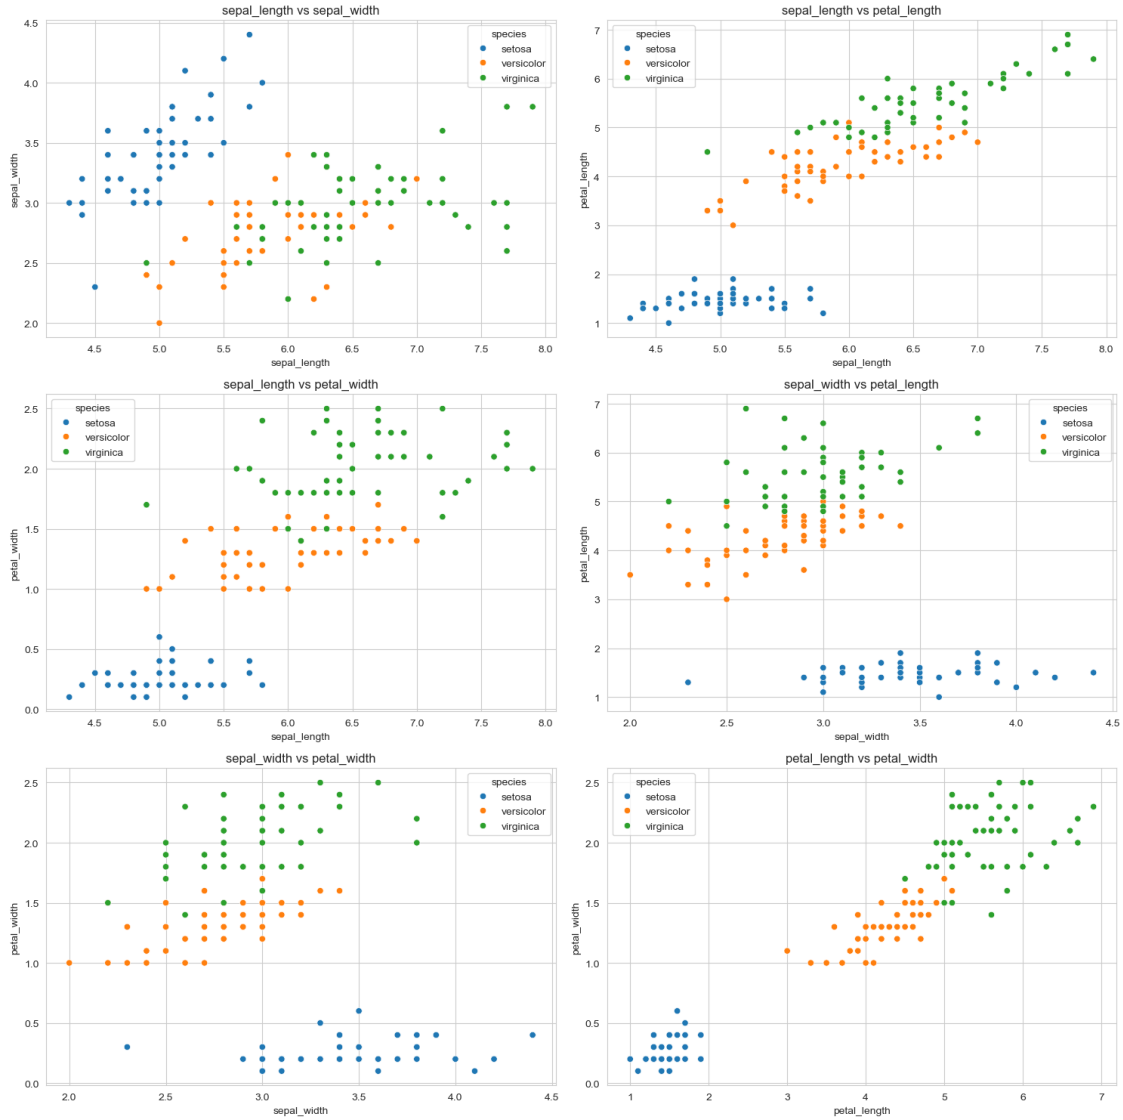
plot_index = 0
for i in range(num_cols_count):
    for j in range(i+1, num_cols_count):
        row = plot_index // 2
        col = plot_index % 2

        # Scatter plot
        sns.scatterplot(x=numerical_columns[i], y=numerical_columns[j],
            ↪hue=output_var, data=iris_df, ax=axes[row, col])
        axes[row, col].set_title(f'{numerical_columns[i]} vs ↪
            ↪{numerical_columns[j]}')

        plot_index += 1

plt.tight_layout()
plt.show()

```



13.3 9.3 Correaltion Numerical Columns


```
[16]: print(iris_df.corr(numeric_only=True))
sns.heatmap(iris_df.corr(numeric_only=True), cmap="YlGnBu", annot=True)
plt.show()
```

	sepal_length	sepal_width	petal_length	petal_width	\
sepal_length	1.000000	-0.117570	0.871754	0.817941	
sepal_width	-0.117570	1.000000	-0.428440	-0.366126	
petal_length	0.871754	-0.428440	1.000000	0.962865	
petal_width	0.817941	-0.366126	0.962865	1.000000	
species_labels	0.782561	-0.426658	0.949035	0.956547	


```

species_labels
sepal_length    0.782561
sepal_width     -0.426658
petal_length    0.949035
petal_width     0.956547
species_labels  1.000000

```



14 10. Data Transformation

15 11. Standization

```

[17]: scaler = StandardScaler()
mean_list = []
std_list = []
for var in ['sepal_length', 'sepal_width', 'petal_length', 'petal_width']:
    mean_list.append(iris_df[var].mean())
    std_list.append(iris_df[var].std())

```

```

print(mean_list)
print(std_list)

# Fit and transform the scaler on the selected columns
scaled_columns = scaler.fit_transform(iris_df[['sepal_length', 'sepal_width',
↪ 'petal_length', 'petal_width']])

# Replace the original columns with the scaled columns
iris_df[['sepal_length_stand', 'sepal_width_stand', 'petal_length_stand',
↪ 'petal_width_stand']] = scaled_columns

print(iris_df)

```

```

[5.8433333333333334, 3.0573333333333337, 3.7580000000000005, 1.1993333333333336]
[0.8280661279778629, 0.435866284936698, 1.7652982332594667, 0.7622376689603465]

```

	sepal_length	sepal_width	petal_length	petal_width	species_labels	\
0	5.1	3.5	1.4	0.2	0	
1	4.9	3.0	1.4	0.2	0	
2	4.7	3.2	1.3	0.2	0	
3	4.6	3.1	1.5	0.2	0	
4	5.0	3.6	1.4	0.2	0	
..	
145	6.7	3.0	5.2	2.3	2	
146	6.3	2.5	5.0	1.9	2	
147	6.5	3.0	5.2	2.0	2	
148	6.2	3.4	5.4	2.3	2	
149	5.9	3.0	5.1	1.8	2	

	species	sepal_length_stand	sepal_width_stand	petal_length_stand	\
0	setosa	-0.900681	1.019004	-1.340227	
1	setosa	-1.143017	-0.131979	-1.340227	
2	setosa	-1.385353	0.328414	-1.397064	
3	setosa	-1.506521	0.098217	-1.283389	
4	setosa	-1.021849	1.249201	-1.340227	
..	
145	virginica	1.038005	-0.131979	0.819596	
146	virginica	0.553333	-1.282963	0.705921	
147	virginica	0.795669	-0.131979	0.819596	
148	virginica	0.432165	0.788808	0.933271	
149	virginica	0.068662	-0.131979	0.762758	

	petal_width_stand
0	-1.315444
1	-1.315444
2	-1.315444
3	-1.315444
4	-1.315444

```

..          ...
145          1.448832
146          0.922303
147          1.053935
148          1.448832
149          0.790671

```

```
[150 rows x 10 columns]
```

16 12. Convert Cat - to - Numerical Columns

17 ML Models

```

[18]: Y=iris_df["species_labels"]
      X=iris_df[['sepal_length_stand', 'sepal_width_stand', 'petal_length_stand',
      ↪ 'petal_width_stand']]
      print(len(Y),len(X))

```

```
150 150
```

```

[19]: X_train, X_test, Y_train, Y_test = train_test_split(X,Y,test_size=0.
      ↪2,random_state = 42)
      print(len(X_train),len(X_test))

```

```
120 30
```

```

[20]: def draw_heatmap(conf_matrix):
      sns.heatmap(conf_matrix, annot=True)
      plt.xlabel('Predicted Labels')
      plt.ylabel('Actual Labels')
      plt.title('Confusion Matrix')
      plt.show()

```

17.1 Logistic Regression

```

[21]: lg_model = LogisticRegression(solver='saga', max_iter=500, random_state=42)
      lg_model.fit(X_train, Y_train)

      print("Model - Logistic Regression")
      score = lg_model.score(X_train, Y_train)
      print('accuracy train score overall :', score)
      score = lg_model.score(X_test, Y_test)
      print('accuracy test score overall :', score)

      y_pred = lg_model.predict(X_test)
      print(classification_report(Y_test, y_pred))
      print(confusion_matrix(Y_test, y_pred))
      conf_matrix = confusion_matrix(Y_test, y_pred)

```

```
draw_heatmap(conf_matrix)
```


Model - Logistic Regression

accuracy train score overall : 0.9666666666666667

accuracy test score overall : 1.0

	precision	recall	f1-score	support
0	1.00	1.00	1.00	10
1	1.00	1.00	1.00	9
2	1.00	1.00	1.00	11
accuracy			1.00	30
macro avg	1.00	1.00	1.00	30
weighted avg	1.00	1.00	1.00	30

```
[[10  0  0]
 [ 0  9  0]
 [ 0  0 11]]
```



17.2 GaussianNB

```
[22]: from sklearn.naive_bayes import GaussianNB, CategoricalNB
      gnb_model = GaussianNB()
      gnb_model.fit(X_train,Y_train)

      print("Model-GaussianNB")
      print("train score",gnb_model.score(X_train,Y_train))
      print("test score",gnb_model.score(X_test,Y_test))

      y_pred = gnb_model.predict(X_test)
      print(classification_report(Y_test, y_pred))
      print(confusion_matrix(Y_test, y_pred))
      conf_matrix = confusion_matrix(Y_test, y_pred)
      draw_heatmap(conf_matrix)
```

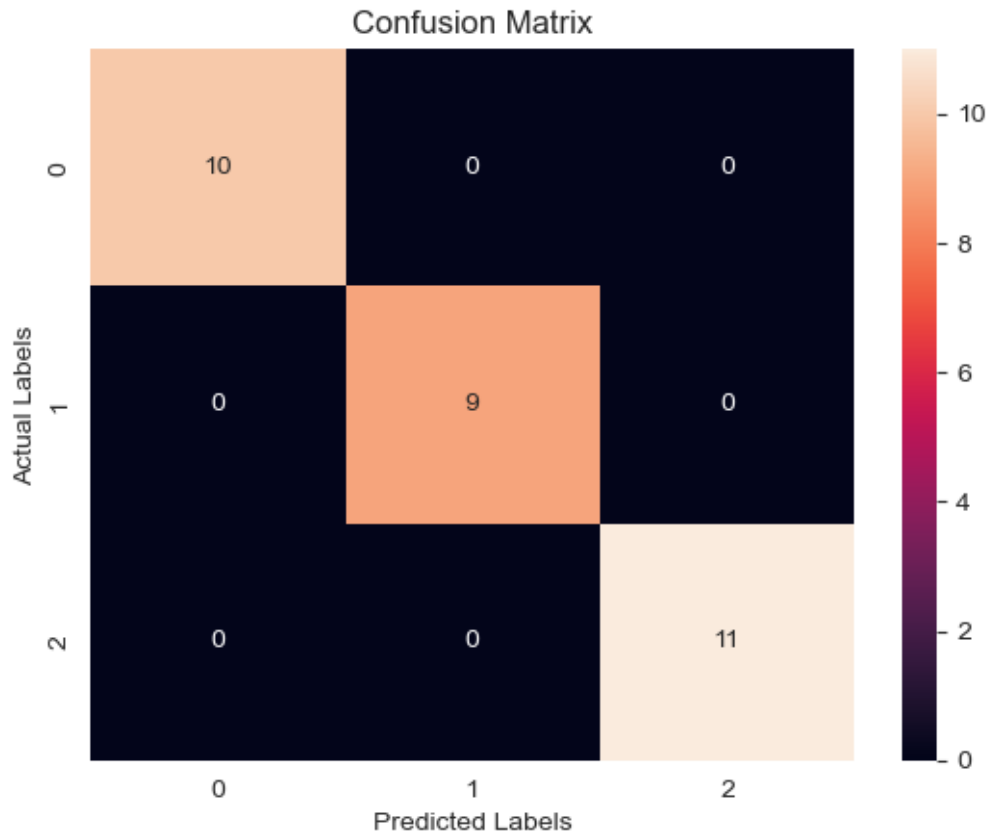
Model-GaussianNB

train score 0.95

test score 1.0

	precision	recall	f1-score	support
0	1.00	1.00	1.00	10
1	1.00	1.00	1.00	9
2	1.00	1.00	1.00	11
accuracy			1.00	30
macro avg	1.00	1.00	1.00	30
weighted avg	1.00	1.00	1.00	30

```
[[10  0  0]
 [ 0  9  0]
 [ 0  0 11]]
```



17.2.1 Save the model

```
[23]: import pickle
pickle.dump(gnb_model, open('iris-model.pkl', 'wb'))
```

18 Support Vector Machine - Classifier

```
[24]: from sklearn.svm import SVC
# Initialize the SVM classifier
svm_linear_classifier = SVC(kernel='linear', random_state=42)

# Train the SVM classifier
svm_linear_classifier.fit(X_train, Y_train)
print("model-Support Vector Machine - kernel - linear -Classifier")

y_pred = svm_linear_classifier.predict(X_train)
# Calculate the accuracy of the model
accuracy = accuracy_score(Y_train, y_pred)
print("Train Accuracy:", accuracy)
```

```

# Predict the classes for test set
y_pred = svm_linear_classifier.predict(X_test)
# Calculate the accuracy of the model
accuracy = accuracy_score(Y_test, y_pred)
print("Test Accuracy:", accuracy)

print(classification_report(Y_test, y_pred))
print(confusion_matrix(Y_test, y_pred))
conf_matrix = confusion_matrix(Y_test, y_pred)
draw_heatmap(conf_matrix)

```

model-Suport Vector Machine - kernel - linear -Classifier

Train Accuracy: 0.9833333333333333


Test Accuracy: 0.9666666666666667

	precision	recall	f1-score	support
0	1.00	1.00	1.00	10
1	1.00	0.89	0.94	9
2	0.92	1.00	0.96	11
accuracy			0.97	30
macro avg	0.97	0.96	0.97	30
weighted avg	0.97	0.97	0.97	30

```

[[10  0  0]
 [ 0  8  1]
 [ 0  0 11]]

```



```
[25]: svm_rbf_classifier = SVC(kernel='rbf', random_state=42)

# Train the SVM classifier
svm_rbf_classifier.fit(X_train, Y_train)
print("model-Suport Vector Machine - Kernel -rbf - Classifier")
y_pred = svm_rbf_classifier.predict(X_train)
# Calculate the accuracy of the model
accuracy = accuracy_score(Y_train, y_pred)
print("Train Accuracy:", accuracy)

# Predict the classes for test set
y_pred = svm_rbf_classifier.predict(X_test)
# Calculate the accuracy of the model
accuracy = accuracy_score(Y_test, y_pred)
print("Test Accuracy:", accuracy)

print(classification_report(Y_test, y_pred))
print(confusion_matrix(Y_test, y_pred))
conf_matrix = confusion_matrix(Y_test, y_pred)
draw_heatmap(conf_matrix)
```



model-Suport Vector Machine - Kernel -rbf - Classifier

Train Accuracy: 0.975

Test Accuracy: 1.0

	precision	recall	f1-score	support
0	1.00	1.00	1.00	10
1	1.00	1.00	1.00	9
2	1.00	1.00	1.00	11
accuracy			1.00	30
macro avg	1.00	1.00	1.00	30
weighted avg	1.00	1.00	1.00	30

```
[[10  0  0]
 [ 0  9  0]
 [ 0  0 11]]
```



```
[26]: svm_poly_classifier = SVC(kernel='poly', random_state=42)

# Train the SVM classifier
```

```

svm_poly_classifier.fit(X_train, Y_train)
print("model-Suport Vector Machine - Kernel -poly - Classifier")

y_pred = svm_poly_classifier.predict(X_train)
# Calculate the accuracy of the model
accuracy = accuracy_score(Y_train, y_pred)
print("Train Accuracy:", accuracy)

# Predict the classes for test set
y_pred = svm_poly_classifier.predict(X_test)

# Calculate the accuracy of the model
accuracy = accuracy_score(Y_test, y_pred)
print("Test Accuracy:", accuracy)

print(classification_report(Y_test, y_pred))
print(confusion_matrix(Y_test, y_pred))
conf_matrix = confusion_matrix(Y_test, y_pred)
draw_heatmap(conf_matrix)

```

model-Suport Vector Machine - Kernel -poly - Classifier

Train Accuracy: 0.9333333333333333


Test Accuracy: 0.9666666666666667

	precision	recall	f1-score	support
0	1.00	1.00	1.00	10
1	0.90	1.00	0.95	9
2	1.00	0.91	0.95	11
accuracy			0.97	30
macro avg	0.97	0.97	0.97	30
weighted avg	0.97	0.97	0.97	30

```

[[10  0  0]
 [ 0  9  0]
 [ 0  1 10]]

```



18.1 Decision Tree

```
[27]: dt_clf = DecisionTreeClassifier(max_leaf_nodes=20,random_state=42)
dt_clf.fit(X_train, Y_train)
print("Model-Decion Tree")

accuracy=dt_clf.score(X_train, Y_train)
print(f"train score: {accuracy}")


accuracy=dt_clf.score(X_test, Y_test)
print(f"test score: {accuracy}")

y_pred=dt_clf.predict(X_test)
print(classification_report(Y_test, y_pred))
print(confusion_matrix(Y_test, y_pred))
conf_matrix = confusion_matrix(Y_test, y_pred)
draw_heatmap(conf_matrix)
```

```
Model-Decion Tree
train score: 1.0
test score: 1.0
```

	precision	recall	f1-score	support
0	1.00	1.00	1.00	10
1	1.00	1.00	1.00	9
2	1.00	1.00	1.00	11
accuracy			1.00	30
macro avg	1.00	1.00	1.00	30
weighted avg	1.00	1.00	1.00	30

```
[[10  0  0]
 [ 0  9  0]
 [ 0  0 11]]
```



18.2 Random Forest

```
[28]: rf_clf= RandomForestClassifier(n_estimators = 1000, random_state = 42,
    ↳max_leaf_nodes=20)
rf_clf.fit(X_train, Y_train)
print("Model- Random Forest Tree")
```

```

accuracy=rf_clf.score(X_train, Y_train)
print(f"train score: {accuracy}")

accuracy=rf_clf.score(X_test, Y_test)
print(f"test score: {accuracy}")

y_pred=rf_clf.predict(X_test)
print(classification_report(Y_test, y_pred))
print(confusion_matrix(Y_test, y_pred))
conf_matrix = confusion_matrix(Y_test, y_pred)
draw_heatmap(conf_matrix)

```

Model- Random Forest Tree

train score: 1.0


test score: 1.0

	precision	recall	f1-score	support
0	1.00	1.00	1.00	10
1	1.00	1.00	1.00	9
2	1.00	1.00	1.00	11
accuracy			1.00	30
macro avg	1.00	1.00	1.00	30
weighted avg	1.00	1.00	1.00	30


```

[[10  0  0]
 [ 0  9  0]
 [ 0  0 11]]

```



18.3 AdaBoost

```
[29]: base_classifier = DecisionTreeClassifier(max_depth=1)
      adaboost_clf = AdaBoostClassifier( n_estimators=50, random_state=42)

      # Train the AdaBoost classifier
      adaboost_clf.fit(X_train, Y_train)

      print("Model-AdaBoost")
      print("train score",adaboost_clf.score(X_train, Y_train))

      # Predict on the test set
      y_pred = adaboost_clf.predict(X_test)

      # Calculate accuracy
      accuracy = accuracy_score(Y_test, y_pred)
      print(f"test score: {accuracy}")

      print(classification_report(Y_test, y_pred))
      print(confusion_matrix(Y_test, y_pred))
```

```
conf_matrix = confusion_matrix(Y_test, y_pred)
draw_heatmap(conf_matrix)
```


Model-AdaBoost

train score 0.9666666666666667

test score: 1.0

	precision	recall	f1-score	support
0	1.00	1.00	1.00	10
1	1.00	1.00	1.00	9
2	1.00	1.00	1.00	11
accuracy			1.00	30
macro avg	1.00	1.00	1.00	30
weighted avg	1.00	1.00	1.00	30

```
[[10  0  0]
 [ 0  9  0]
 [ 0  0 11]]
```



18.4 GradientBoostingClassifier

```
[30]: gdb_clf = GradientBoostingClassifier(n_estimators=20, learning_rate=0.1,
      ↪max_depth=1, random_state=42)
gdb_clf.fit(X_train, Y_train)
print("model-Gradient Boosting Classifier")

accuracy = gdb_clf.score(X_train, Y_train)
print("Train Accuracy:", accuracy)

accuracy = gdb_clf.score(X_test, Y_test)
print("Test Accuracy:", accuracy)

print(classification_report(Y_test, y_pred))
print(confusion_matrix(Y_test, y_pred))
conf_matrix = confusion_matrix(Y_test, y_pred)
draw_heatmap(conf_matrix)
```


model-Gradient Boosting Classifier

Train Accuracy: 0.9583333333333334

Test Accuracy: 1.0

	precision	recall	f1-score	support
0	1.00	1.00	1.00	10
1	1.00	1.00	1.00	9
2	1.00	1.00	1.00	11
accuracy			1.00	30
macro avg	1.00	1.00	1.00	30
weighted avg	1.00	1.00	1.00	30

```
[[10  0  0]
 [ 0  9  0]
 [ 0  0 11]]
```

18.5 XGBClassifier

```
[31]: from xgboost import XGBClassifier
xgmodel = XGBClassifier()
xgmodel.fit(X_train, Y_train)

print("model- XGB Classifier")
# Make predictions on the test set
y_pred = xgmodel.predict(X_train)
accuracy = accuracy_score(Y_train, y_pred)
print("Test Accuracy:", accuracy)
# Evaluate the model

# Make predictions on the test set
y_pred = xgmodel.predict(X_test)
accuracy = accuracy_score(Y_test, y_pred)
print("Test Accuracy:", accuracy)

print(classification_report(Y_test, y_pred))
```

```
print(confusion_matrix(Y_test, y_pred))
conf_matrix = confusion_matrix(Y_test, y_pred)
draw_heatmap(conf_matrix)
```


model- XGB Classifier

Test Accuracy: 1.0

Test Accuracy: 1.0

	precision	recall	f1-score	support
0	1.00	1.00	1.00	10
1	1.00	1.00	1.00	9
2	1.00	1.00	1.00	11
accuracy			1.00	30
macro avg	1.00	1.00	1.00	30
weighted avg	1.00	1.00	1.00	30

```
[[10  0  0]
 [ 0  9  0]
 [ 0  0 11]]
```



[]: